

**5<sup>th</sup>**  
**Week**

# 다섯 번째 뵙겠습니다 ?!

▷ 잠시만 기다렸다가 30분 되면 시작하겠습니다~^^

▷ 지난 주에는 조금 일찍 끝났으니, 이번 주에는 ... ?!

- 신나게 달려봅시다~~~!!!

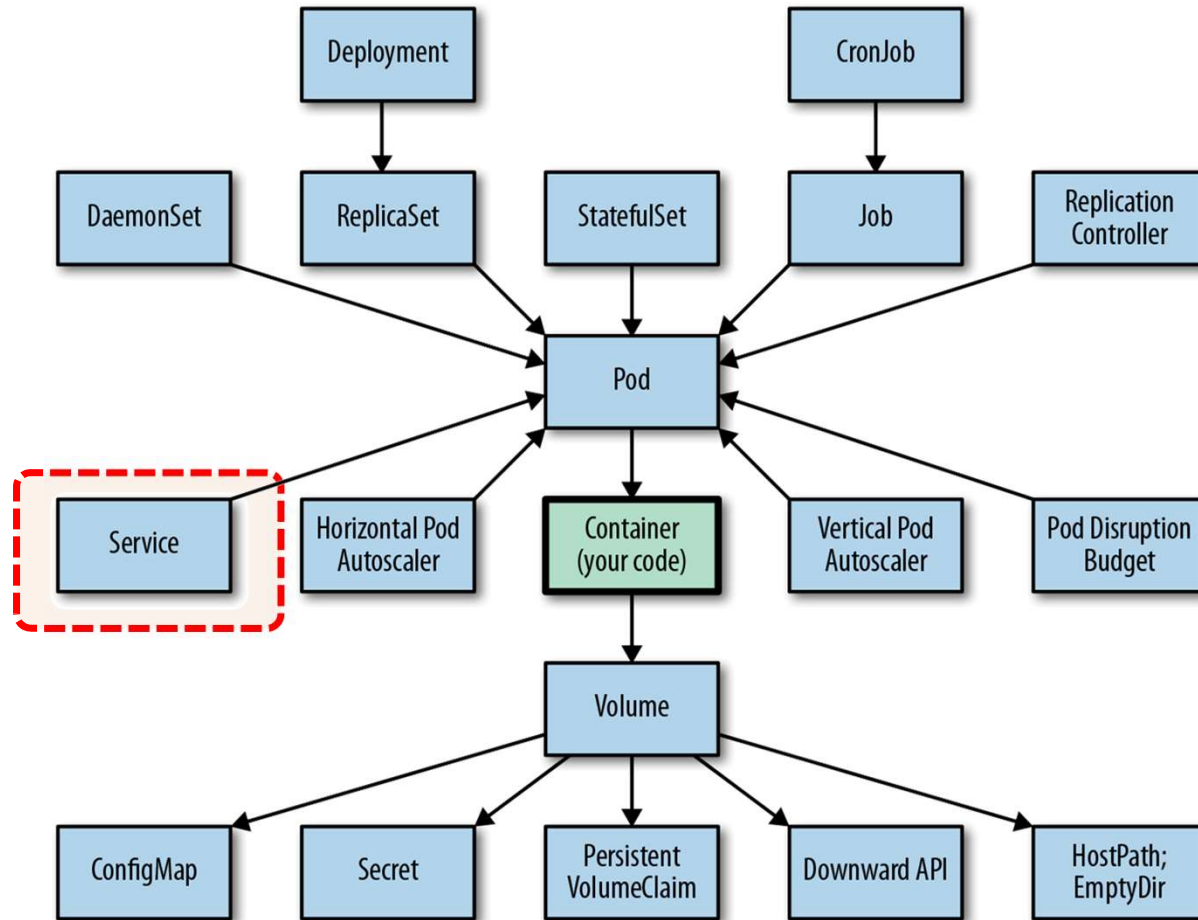
▷ Camera는 가급적 켜 주시면 대단히 감사하겠습니다 !!!

- 너무 부끄러우면 Snap Camera를 사용하시는 것까지는~ ^^

▷ 오늘 수업 자료는 아래 링크에서 다운로드 받으실 수 있어요.

- <https://github.com/whatwant-school/kubernetes>

## Service



※ 참고 : <https://www.oreilly.com/library/view/kubernetes-patterns/9781492050278/ch01.html>



# 지난 수업 복습

**<https://kahoot.it/>**



# **Service**

**Ingress / LoadBalancer**





# **Flip Learning**

**(Service - Ingress / LoadBalancer)**

**권재혁 님**



# **Kubernetes**

## **Network - Ingress**

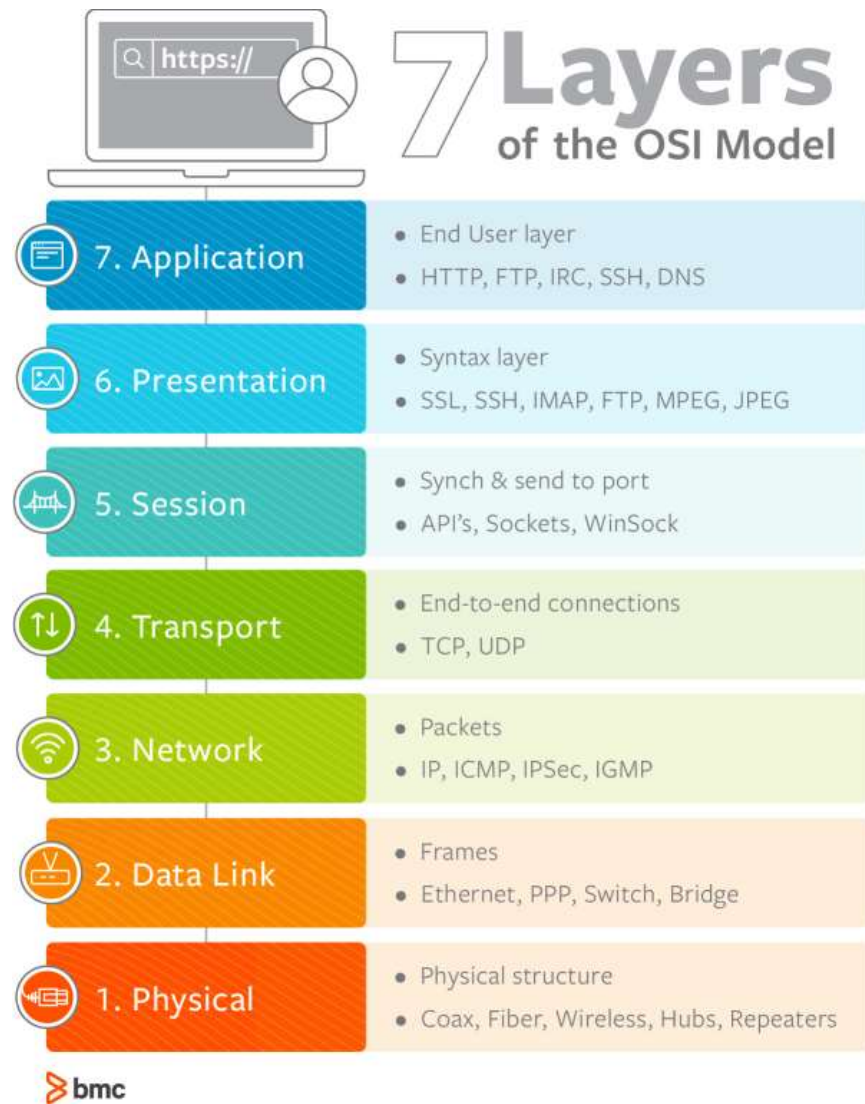
# Ingress Controller

- Ingress는 cluster 외부에서 내부 Service로 HTTP와 HTTPS 경로 노출
- 클라이언트가 요청한 호스트와 경로에 따라 요청을 전달할 서비스가 결정
  - . HTTP(S)기반의 L7 LoadBalancing 기능을 제공하는 컴포넌트
- Ingress 구현체 : 각 구현체마다 설정 방법 차이 존재
  - . 구글 클라우드 : <https://github.com/kubernetes/ingress-gce>
  - . 오픈소스 (Ingress NGINX Controller) : <https://github.com/kubernetes/ingress-nginx>
  - . 오픈소스 (NGINX Ingress Controller) : <https://github.com/nginxinc/kubernetes-ingress>
  - . 상용 (F5 BIG IP Controller) : <http://clouddocs.f5.com/products/connectors/k8s-bigip-ctrlr>



※ 참고 : <https://kubernetes.io/ko/docs/concepts/services-networking/ingress/>

# OSI 7 Layers



※ 참고 : <https://www.bmc.com/blogs/osi-model-7-layers/>

# Ingress Controller Install

Kubernetes에서 배포하고 있는 패키지로 설치해보는데, 아래 링크에서 bare-metal 설치 방법을 찾아보면 된다

- <https://github.com/kubernetes/ingress-nginx/blob/main/docs/deploy/index.md#bare-metal-clusters>

```
remote > kubectl apply -f https://raw.githubusercontent.com/kubernetes/ingress-nginx/controller-v1.3.0/deploy/static/provider/baremetal/deploy.yaml

namespace/ingress-nginx created
serviceaccount/ingress-nginx created
serviceaccount/ingress-nginx-admission created
role.rbac.authorization.k8s.io/ingress-nginx created
role.rbac.authorization.k8s.io/ingress-nginx-admission created
clusterrole.rbac.authorization.k8s.io/ingress-nginx created
clusterrole.rbac.authorization.k8s.io/ingress-nginx-admission created
rolebinding.rbac.authorization.k8s.io/ingress-nginx created
rolebinding.rbac.authorization.k8s.io/ingress-nginx-admission created
clusterrolebinding.rbac.authorization.k8s.io/ingress-nginx created
clusterrolebinding.rbac.authorization.k8s.io/ingress-nginx-admission created
configmap/ingress-nginx-controller created
service/ingress-nginx-controller created
service/ingress-nginx-controller-admission created
deployment.apps/ingress-nginx-controller created
job.batch/ingress-nginx-admission-create created
job.batch/ingress-nginx-admission-patch created
ingressclass.networking.k8s.io/nginx created
validatingwebhookconfiguration.admissionregistration.k8s.io/ingress-nginx-admission created
```

# Ingress Controller Check

```
remote > kubectl get namespaces
```

NAME	STATUS	AGE
default	Active	23d
ingress-nginx	Active	2m19s
...		

```
remote > kubectl get pods -o wide --namespace kube-system
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
...								
nginx-proxy-worker1	1/1	Running	11 (5m42s ago)	23d	192.168.100.201	worker1	<none>	<none>
nginx-proxy-worker2	1/1	Running	9 (5m34s ago)	23d	192.168.100.202	worker2	<none>	<none>
...								

```
remote > kubectl get pods -o wide --namespace ingress-nginx
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
ingress-nginx-admission-create-6j9cv	0/1	Completed	0	3m49s	10.233.103.75	worker2	<none>	<none>
ingress-nginx-admission-patch-qmwvf	0/1	Completed	1	3m49s	10.233.110.24	worker1	<none>	<none>
ingress-nginx-controller-77cb5dbf4d-k5qmr	1/1	Running	0	3m49s	10.233.110.25	worker1	<none>	<none>

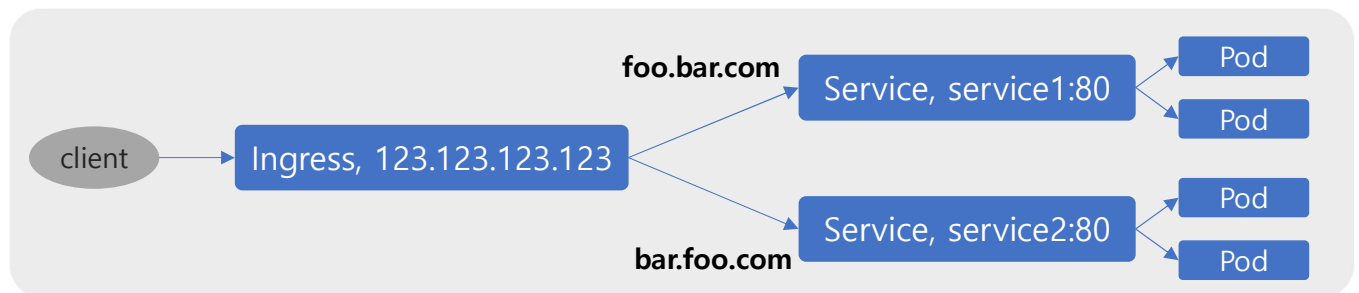
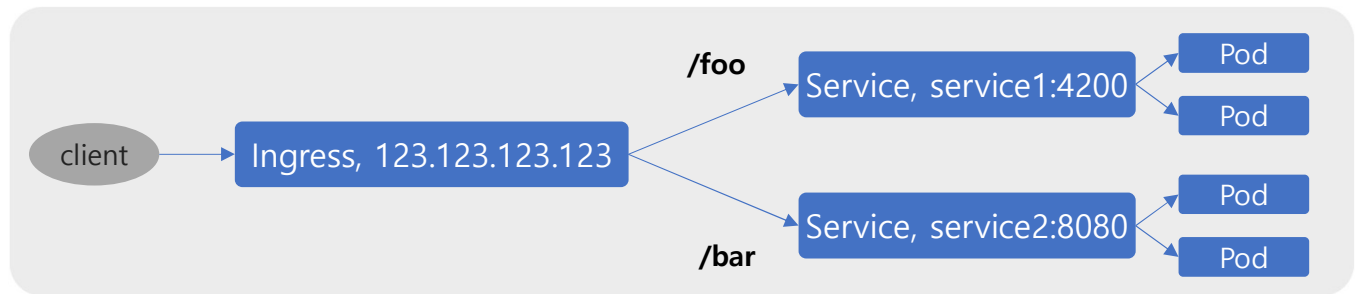
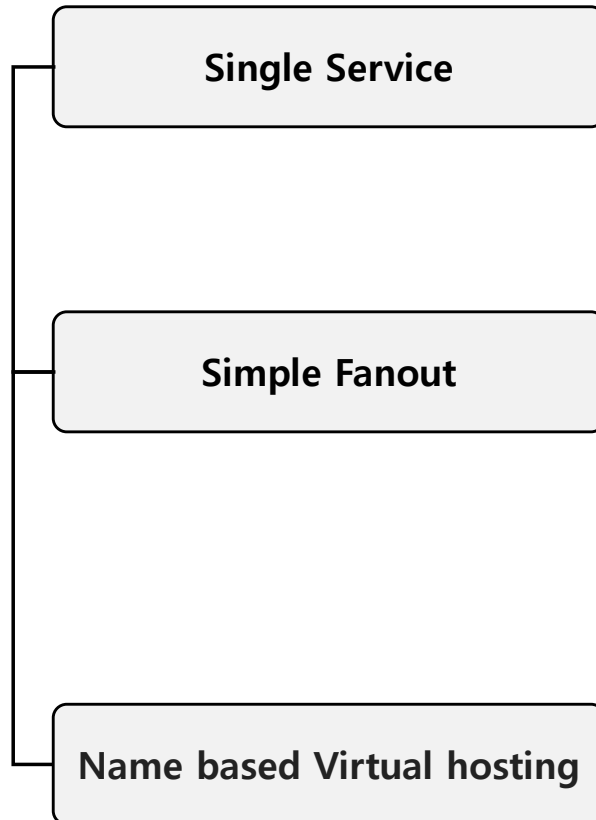
```
remote > kubectl get serviceaccounts -o wide --namespace ingress-nginx
```

NAME	SECRETS	AGE
default	1	4m28s
ingress-nginx	1	4m28s
ingress-nginx-admission	1	4m28s





# Types of Ingress



※ 참고 : <https://kubernetes.io/docs/concepts/services-networking/ingress/>



# Ready: container images

별개의 사이트 구분을 할 수 있도록 웹서비스를 2종 준비했다.

기존 `node-web:1.0`과 구분이 될 정도만 살짝 변경한 후 `node-web:2.0` 버전으로 DockerHub에 push 했다.

```
const http = require('http');
const os = require('os');

console.log("node-web server starting...");

var handler = function(request, response) {
  console.log("Received request from " + request.connection.remoteAddress);
  response.writeHead(200);
  response.end("You've hit " + os.hostname() + "(Ver2.0)\n");
};

var www = http.createServer(handler);
www.listen(8080);
```

**app.js**

```
FROM node:latest
ADD app.js /app.js
ENTRYPOINT ["node", "app.js"]
```

**Dockerfile**

DockerHub에 이미 있는 만들어 놓은  
이미지를 사용해도 좋다.

```
remote > git clone https://github.com/whatwant-school/kubernetes.git
remote > cd kubernetes/05-Ingress-LoadBalancer/hands-on

remote > docker build -t node-web:2.0 .

remote > docker tag node-web:2.0 <user-id>/node-web:2.0

remote > docker push <user-id>/node-web:2.0
```

# Ready: ReplicaSet YAML

웹사이트 2종을 각각 띄울 `ReplicaSet`을 준비하자.

```
apiVersion: apps/v1
kind: ReplicaSet

metadata:
  name: rs-web-v1

spec:
  replicas: 3
  selector:
    matchLabels:
      app: node-web-v1

  template:
    metadata:
      labels:
        app: node-web-v1

    spec:
      containers:
        - name: node-web
          image: whatwant/node-web:1.0
          ports:
            - containerPort: 8080
          imagePullPolicy: Always
```

**rs-web-v1.yaml**

```
apiVersion: apps/v1
kind: ReplicaSet

metadata:
  name: rs-web-v2

spec:
  replicas: 3
  selector:
    matchLabels:
      app: node-web-v2

  template:
    metadata:
      labels:
        app: node-web-v2

    spec:
      containers:
        - name: node-web
          image: whatwant/node-web:2.0
          ports:
            - containerPort: 8080
          imagePullPolicy: Always
```

**rs-web-v2.yaml**

# Ready: Create ReplicaSet

앞에서 작성해 놓은 ReplicaSet을 이용해서 Pods를 생성하자.

```
remote > git clone https://github.com/whatwant-school/kubernetes.git
remote > cd kubernetes/05-Ingress-LoadBalancer/hands-on
```

```
remote > kubectl create -f rs-web-v1.yaml
```

```
replicaset.apps/rs-web-v1 created
```

```
remote > kubectl create -f rs-web-v2.yaml
```

```
replicaset.apps/rs-web-v2 created
```

```
remote > kubectl get pods -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
rs-web-v1-ghbqs	1/1	Running	0	2m25s	10.233.110.29	worker1	<none>	<none>
rs-web-v1-m6gtz	1/1	Running	0	2m25s	10.233.103.79	worker2	<none>	<none>
rs-web-v1-rxpxz	1/1	Running	0	2m25s	10.233.103.80	worker2	<none>	<none>
rs-web-v2-cmvl2	1/1	Running	0	2m22s	10.233.103.82	worker2	<none>	<none>
rs-web-v2-k6src	1/1	Running	0	2m22s	10.233.110.30	worker1	<none>	<none>
rs-web-v2-qsh68	1/1	Running	0	2m22s	10.233.103.81	worker2	<none>	<none>

# Ready: Create Service

`ClusterIP` 유형으로 해도 되지만, 여기에서는 `NodePort` 유형으로 `Service`를 만들어보자.

```
apiVersion: v1
kind: Service
metadata:
  name: svc-node-web-v1
spec:
  type: NodePort
  ports:
    - port: 80
      targetPort: 8080
      nodePort: 30001
  selector:
    app: node-web-v1
```

**svc-node-web-nodeport-v1.yaml**

```
apiVersion: v1
kind: Service
metadata:
  name: svc-node-web-v2
spec:
  type: NodePort
  ports:
    - port: 80
      targetPort: 8080
      nodePort: 30002
  selector:
    app: node-web-v2
```

**svc-node-web-nodeport-v2.yaml**

```
remote > cd kubernetes/05-Ingress-LoadBalancer/hands-on
```

```
remote > kubectl create -f svc-node-web-nodeport-v1.yaml
```

```
remote > kubectl create -f svc-node-web-nodeport-v2.yaml
```

```
remote > kubectl get services -o wide
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE	SELECTOR
kubernetes	ClusterIP	10.233.0.1	<none>	443/TCP	24d	<none>
svc-node-web-v1	NodePort	10.233.24.241	<none>	80:30001/TCP	11s	app=node-web-v1
svc-node-web-v2	NodePort	10.233.26.196	<none>	80:30002/TCP	7s	app=node-web-v2





# Status

```
remote > kubectl get replicaset -o wide
```

NAME	DESIRED	CURRENT	READY	AGE	CONTAINERS	IMAGES	SELECTOR
rs-web-v1	3	3	3	8m55s	node-web	whatwant/node-web:1.0	app=node-web-v1
rs-web-v2	3	3	3	8m52s	node-web	whatwant/node-web:2.0	app=node-web-v2

```
remote > kubectl get pods -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED	NODE	READINESS	GATES
rs-web-v1-ghbqs	1/1	Running	0	9m24s	10.233.110.29	worker1	<none>		<none>	
rs-web-v1-m6gtz	1/1	Running	0	9m24s	10.233.103.79	worker2	<none>		<none>	
rs-web-v1-rxpxz	1/1	Running	0	9m24s	10.233.103.80	worker2	<none>		<none>	
rs-web-v2-cmvl2	1/1	Running	0	9m21s	10.233.103.82	worker2	<none>		<none>	
rs-web-v2-k6src	1/1	Running	0	9m21s	10.233.110.30	worker1	<none>		<none>	
rs-web-v2-qsh68	1/1	Running	0	9m21s	10.233.103.81	worker2	<none>		<none>	

```
remote > kubectl get services -o wide
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE	SELECTOR
kubernetes	ClusterIP	10.233.0.1	<none>	443/TCP	24d	<none>
svc-node-web-v1	NodePort	10.233.24.241	<none>	80:30001/TCP	11s	app=node-web-v1
svc-node-web-v2	NodePort	10.233.26.196	<none>	80:30002/TCP	7s	app=node-web-v2

```
remote > curl -s http://192.168.100.200:30001
```

You've hit rs-web-v1-m6gtz

```
remote > curl -s http://192.168.100.200:30002
```

You've hit rs-web-v2-k6src (Ver2.0)

# Create Ingress (Simple Fanout)

## ingress-node-web.yaml

```
apiVersion: networking.k8s.io/v1
kind: Ingress
```

### metadata:

```
  name: ingress-node-web
```

### annotations:

```
  kubernetes.io/ingress.class: "nginx"
```

```
  nginx.ingress.kubernetes.io/rewrite-target: /
```

### spec:

#### rules:

##### - http:

#### paths:

```
  - path: /foo
```

```
    pathType: Prefix
```

#### backend:

##### service:

```
  name: svc-node-web-v1
```

##### port:

```
  number: 80
```

```
  - path: /bar
```

```
    pathType: Prefix
```

#### backend:

##### service:

```
  name: svc-node-web-v2
```

##### port:

```
  number: 80
```

```
remote > cd kubernetes/05-Ingress-LoadBalancer/hands-on
```

```
remote > kubectl create -f ingress-node-web.yaml
```

```
ingress.networking.k8s.io/ingress-node-web created
```

```
remote > kubectl get ingresses -o wide
```

NAME	CLASS	HOSTS	ADDRESS	PORTS	AGE
ingress-node-web	<none>	*	192.168.100.201	80	11s

```
remote > kubectl get services --namespace ingress-nginx
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
ingress-nginx-controller	NodePort	10.233.63.114	<none>	80:30977/TCP, 443:30168/TCP	13h
ingress-nginx-controller-admission	ClusterIP	10.233.59.2	<none>	443/TCP	13h

```
remote > curl -s http://192.168.100.201:30977/foo
```

```
You've hit rs-web-v1-m6gtz
```

```
remote > curl -s http://192.168.100.201:30977/bar
```

```
You've hit rs-web-v2-cmv12 (Ver2.0)
```

# Describe Ingress

```
remote > kubectl describe ingress ingress-node-web
```


```
Name:          ingress-node-web
Labels:        <none>
Namespace:     default
Address:       192.168.100.201
Default backend: default-http-backend:80 (<error: endpoints "default-http-backend" not found>)
```

```
Rules:
  Host      Path      Backends
  ----      -
  *
    /foo    svc-node-web-v1:80 (10.233.103.79:8080,10.233.103.80:8080,10.233.110.29:8080)
    /bar    svc-node-web-v2:80 (10.233.103.81:8080,10.233.103.82:8080,10.233.110.30:8080)
```

```
Annotations:  kubernetes.io/ingress.class: nginx
              nginx.ingress.kubernetes.io/rewrite-target: /
```

```
Events:
  Type    Reason    Age           From          Message
  ----    -
  Normal  Sync      12m (x2 over 13m)  nginx-ingress-controller  Scheduled for sync
```

# Design Thinking





# externalTrafficPolicy

NodePort  
공부하고 이어서 바로 했으면  
더 좋았을 것 같다.

# Ready

- 앞에서 NodePort 실습했던 내용을 활용해서 살펴보도록 하겠다.

```
remote > cd kubernetes/03-RS-DS-JOB-CJ/hands-on
```

```
remote > kubectl create -f rs-web.yaml
```

```
remote > cd kubernetes/04-ClusterIP-NodePort-ExternalName/hands-on
```

```
remote > kubectl create -f svc-web-node.yaml
```

```
remote > kubectl get pods -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED	NODE	READINESS	GATES
rs-web-d7sd5	1/1	Running	0	2m33s	10.233.110.31	worker1	<none>		<none>	
rs-web-h6cpk	1/1	Running	0	2m33s	10.233.103.85	worker2	<none>		<none>	
rs-web-ltscf	1/1	Running	0	2m33s	10.233.103.84	worker2	<none>		<none>	

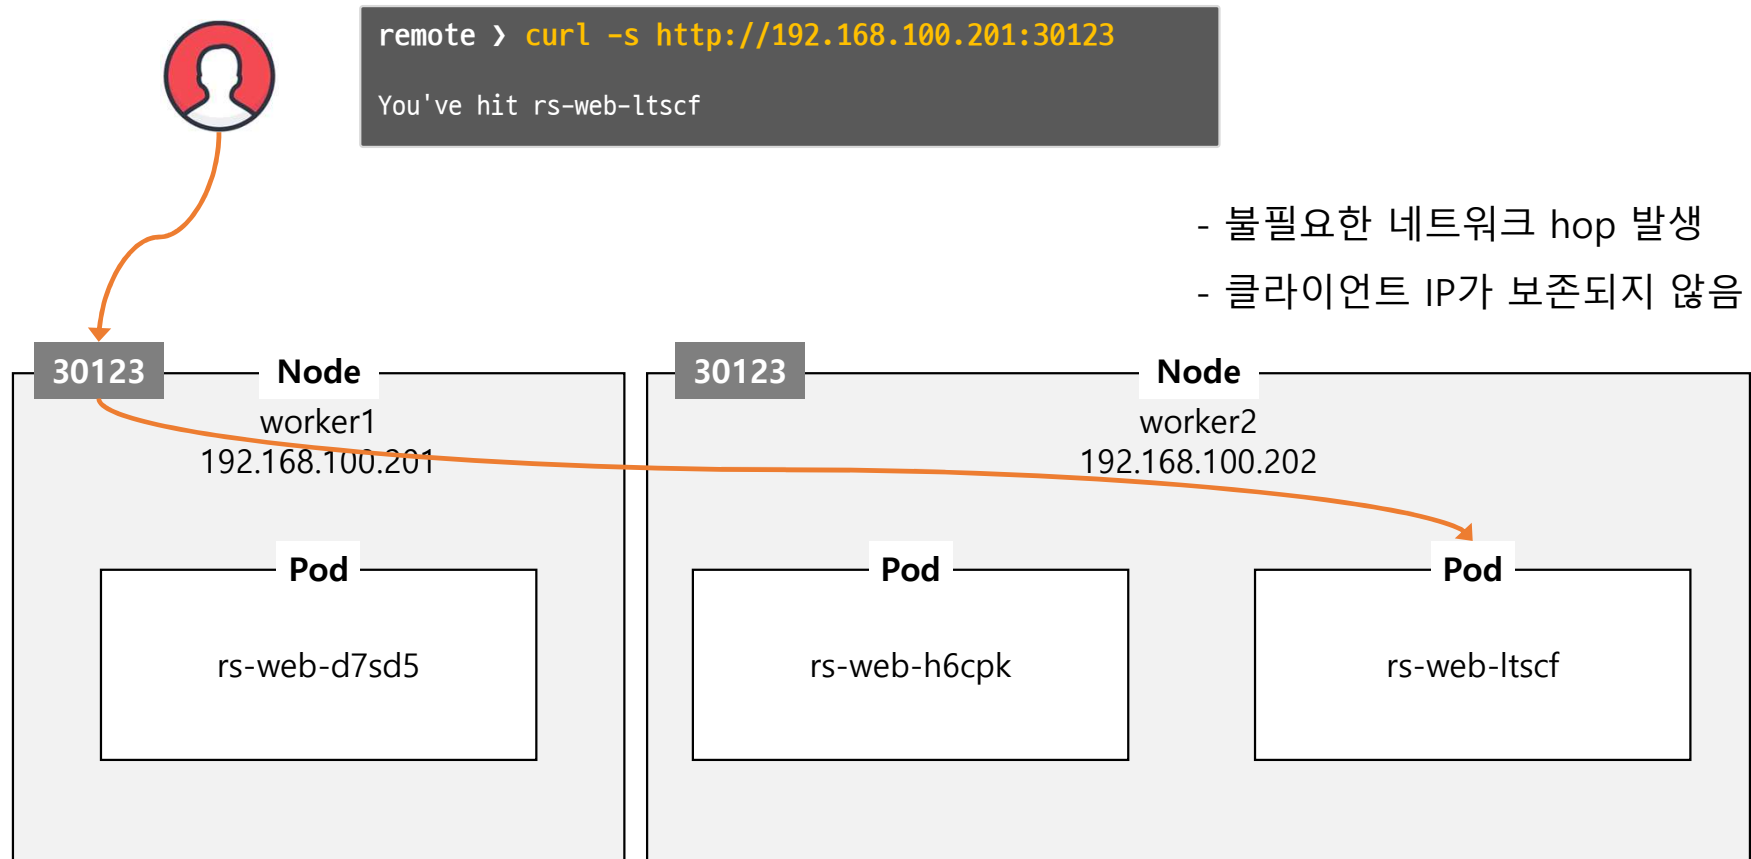
```
remote > kubectl get services -o wide
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE	SELECTOR
kubernetes	ClusterIP	10.233.0.1	<none>	443/TCP	24d	<none>
svc-web-node	NodePort	10.233.7.149	<none>	80:30123/TCP	77s	app=node-web

```
remote > curl -s http://192.168.100.201:30123
```

```
You've hit rs-web-ltscf
```

# Problem





# externalTrafficPolicy

## svc-web-node-et.yaml

```
apiVersion: v1
kind: Service

metadata:
  name: svc-web-node-et

spec:
  type: NodePort

  ports:
    - port: 80
      targetPort: 8080
      nodePort: 30125

  selector:
    app: node-web

externalTrafficPolicy: Local
```

```
remote > cd kubernetes/05-Ingress-LoadBalancer/hands-on
```

```
remote > kubectl create -f svc-web-node-et.yaml
```

```
service/svc-web-node-et created
```


```
remote > kubectl get services -o wide
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE	SELECTOR
kubernetes	ClusterIP	10.233.0.1	<none>	443/TCP	24d	<none>
svc-web-node	NodePort	10.233.7.149	<none>	80:30123/TCP	89m	app=node-web
svc-web-node-et	NodePort	10.233.6.235	<none>	80:30125/TCP	27s	app=node-web

```
remote > curl -s http://192.168.100.201:30125
```

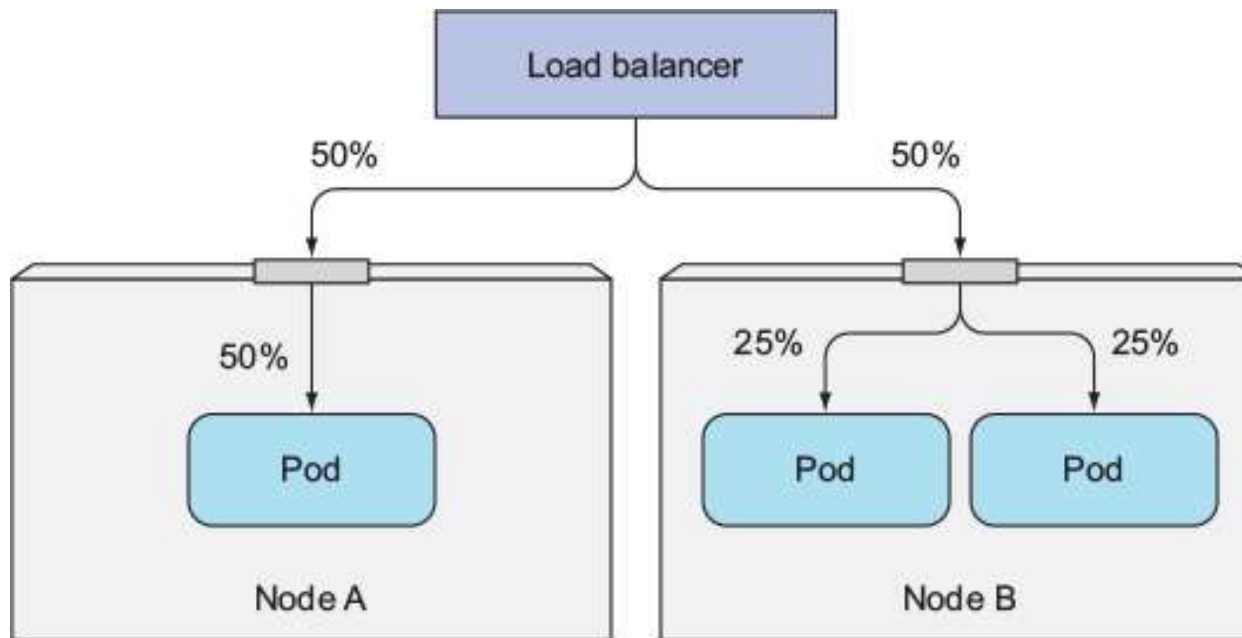
```
You've hit rs-web-d7sd5
```

- externalTrafficPolicy 기본값은 Cluster



# But,

- LoadBalancer를 사용하면서 externalTrafficPolicy를 적용하게 되면, 오히려 균등 배부가 되지 않을 수도 있다.



※ "externalTrafficPolicy: Local" 설정에 따른 문제점 : 균등히 배부되지 않을 수 있음

※ 참고 : <https://livebook.manning.com/book/kubernetes-in-action/chapter-5/230>



**Break**

**돌아오셨으면 채팅창에  
복귀!  
타이핑하기!**


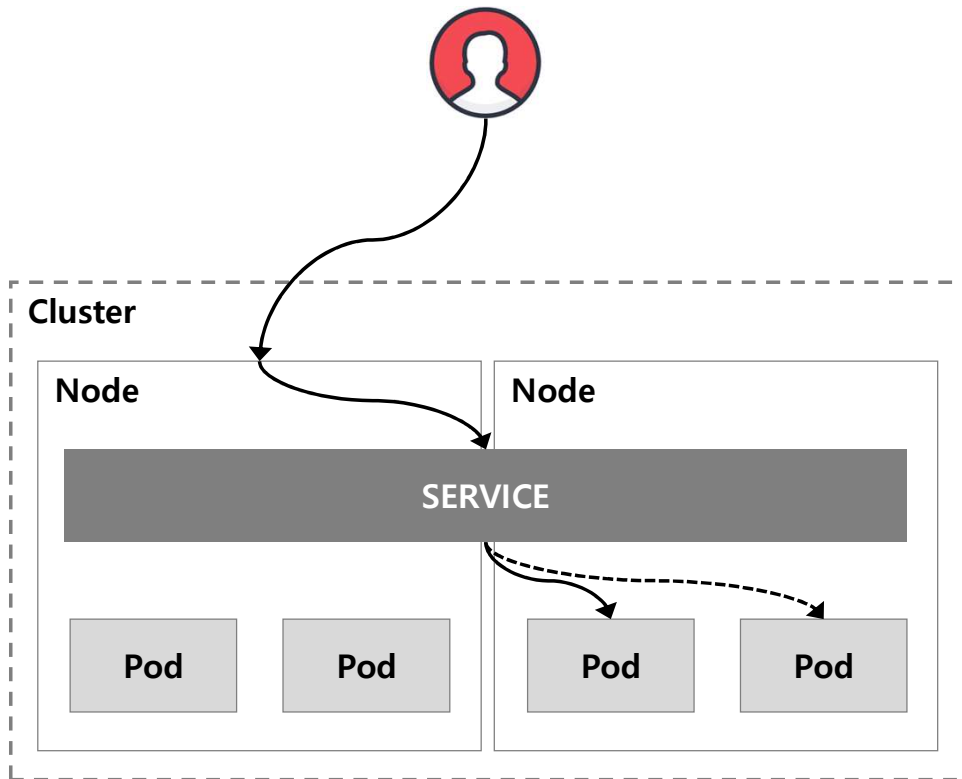


# **Kubernetes**

## **Network - LoadBalancer**

# Why LoadBalancer ?

## NodePort vs LoadBalancer







# **MetalLB**

# **Installation**

# Requirement - 1/3

- MetalLB requires the following to function:

- ① A Kubernetes cluster, running Kubernetes **1.13.0 or later**, that does not already have network load-balancing functionality.
- ② A cluster network configuration that can **coexist with MetalLB**.
- ③ Some **IPv4 addresses** for MetalLB to hand out.
- ④ When using the BGP operating mode, you will need one or more routers capable of speaking BGP.
- ⑤ Traffic on **port 7946** (TCP & UDP) must be allowed between nodes, as required by memberlist.

## ① Kubernetes version

```
remote > kubectl get nodes -o wide
```

NAME	STATUS	ROLES	AGE	VERSION	INTERNAL-IP	EXTERNAL-IP	OS-IMAGE	KERNEL-VERSION	CONTAINER-RUNTIME
master	Ready	control-plane,master	24d	v1.23.7	192.168.100.200	<none>	Ubuntu 20.04.4 LTS	5.4.0-122-generic	containerd://1.6.4
worker1	Ready	<none>	24d	v1.23.7	192.168.100.201	<none>	Ubuntu 20.04.4 LTS	5.4.0-124-generic	containerd://1.6.4
worker2	Ready	<none>	24d	v1.23.7	192.168.100.202	<none>	Ubuntu 20.04.4 LTS	5.4.0-124-generic	containerd://1.6.4

Kubernetes 1.13.0 버전 이상이고, 다른 load-balancing 기능을 갖고 있지 않으면 된다.

※ 참고 : <https://metallb.universe.tf/requirements>

# Requirement - 2/3

## ② Cluster Network

Calico 환경에서는 알려진 문제가 있지만,  
BGP를 사용할 경우에만 해당하는 문제이니 일단 무시하고 진행~

```
remote > kubectl get pods --namespace kube-system
```

NAME	READY	STATUS	RESTARTS	AGE
calico-kube-controllers-6dd874f784-rjxtz	1/1	Running	14 (4h11m ago)	24d
calico-node-599ck	1/1	Running	10 (4h12m ago)	24d
calico-node-qlhvf	1/1	Running	10 (4h11m ago)	24d
calico-node-tpwvg	1/1	Running	20 (4h11m ago)	24d
coredns-76b4fb4578-kc7vm	1/1	Running	10 (4h12m ago)	24d
coredns-76b4fb4578-zbtvb	1/1	Running	9 (4h11m ago)	24d
dns-autoscaler-7979fb6659-p5fpm	1/1	Running	10 (4h12m ago)	24d
kube-apiserver-master	1/1	Running	11 (4h12m ago)	24d
kube-controller-manager-master	1/1	Running	11 (4h12m ago)	24d
...				

Network addon	Compatible
Antrea	Yes (Tested on version 1.4 and 1.5)
Calico	Mostly (see known issues)
Canal	Yes
Cilium	Yes
Flannel	Yes
Kube-ovn	Yes
Kube-router	Mostly (see known issues)
Weave Net	Mostly (see known issues)

## ③ IPv4 addresses

**DHCP 서버 설정**

공유기의 DHCP 서버 기능을 관리합니다.

DHCP 서버 사용함 : ☒

DHCP IP주소 범위 :  to  (1 ~ 254 사이의 숫자)


MetalLB에 할당할 수 있는 IP들을 확보해야 한다.

※ 참고 : <https://metallb.universe.tf/installation/network-addons/>

로컬 클라이언트 상태		
MAC 주소	IP 주소	장치명
08:00:27:00:00:01	192.168.100.102	vm-k8s-master-01
08:00:27:00:00:02	192.168.100.107	vm-k8s-master-02
08:00:27:00:00:03	192.168.100.101	vm-k8s-master-03
08:00:27:00:00:04	192.168.100.200	vm-k8s-master-04
08:00:27:00:00:05	192.168.100.201	vm-k8s-master-05
08:00:27:00:00:06	192.168.100.202	vm-k8s-master-06

## Requirement - 3/3

### ④ When using the BGP(Boarder Gateway Protocol)...



MetalLB는 L2/L3 2가지 방식 중 하나를 선택해서 설치할 수 있다. L3(BGP) 방식을 선택할 경우 추가적으로 고려할 사항들이 있다. 하지만, 여기에서는 일단 L2로 진행할 계획이기에 Skip ...

<b>L2</b>	Data Link	Mac 주소 기반	
<b>L3</b>	Network	IP 주소 기반	
<b>L4</b>	Transport	Port 기반	TCP, UDP
<b>L7</b>	Application	요청(URL) 기반	HTTP, HTTPS

### ⑤ Traffic on port 7946 (TCP & UDP)

MetalLB가 동작하기 위해서는 Node끼리 port 7946 (TCP/UDP) 통신이 되어야 한다.

보통의 상황에서는 문제가 없겠지만, 방화벽이 있는 환경에서는 해당 port가 열려 있도록 주의해야 한다.

# Preparation

"If you're using kube-proxy in **IPVS** mode, since Kubernetes v1.14.2 you have to enable strict ARP mode"

- 실습환경 구축을 가이드한대로 진행했다고 하면, kube-proxy 기본 설정이 `iptables` 일 것이다.
- 만약, `IPVS` 설정으로 되어 있다면 strictARP 설정을 enable로 설정해줘야 한다.
- kube-proxy 설정 내역을 확인해 보려면 다음과 같이 해보면 된다.

```
remote > kubectl describe configmap kube-proxy --namespace kube-system
```

```
...
iptables:
  masqueradeAll: false
  masqueradeBit: 14
  minSyncPeriod: 0s
  syncPeriod: 30s
ipvs:
  excludeCIDRs: []
  minSyncPeriod: 0s
  scheduler: rr
  strictARP: false
  syncPeriod: 30s
  tcpFinTimeout: 0s
  tcpTimeout: 0s
  udpTimeout: 0s
kind: KubeProxyConfiguration
metricsBindAddress: 127.0.0.1:10249
mode: iptables
...
```

※ 참고 : <https://metallb.universe.tf/installation/>



# Install MetalLB

```
remote > kubectl apply -f https://raw.githubusercontent.com/metallb/metallb/v0.13.4/config/manifests/metallb-native.yaml
```

```
namespace/metallb-system created
customresourcedefinition.apiextensions.k8s.io/addresspools.metallb.io created
customresourcedefinition.apiextensions.k8s.io/bfdprofiles.metallb.io created
customresourcedefinition.apiextensions.k8s.io/bgpadvertisements.metallb.io created
customresourcedefinition.apiextensions.k8s.io/bgppeers.metallb.io created
customresourcedefinition.apiextensions.k8s.io/communities.metallb.io created
customresourcedefinition.apiextensions.k8s.io/ipaddresspools.metallb.io created
customresourcedefinition.apiextensions.k8s.io/l2advertisements.metallb.io created
serviceaccount/controller created
serviceaccount/speaker created
Warning: policy/v1beta1 PodSecurityPolicy is deprecated in v1.21+, unavailable in v1.25+
podsecuritypolicy.policy/controller created
podsecuritypolicy.policy/speaker created
role.rbac.authorization.k8s.io/controller created
role.rbac.authorization.k8s.io/pod-lister created
clusterrole.rbac.authorization.k8s.io/metallb-system:controller created
clusterrole.rbac.authorization.k8s.io/metallb-system:speaker created
rolebinding.rbac.authorization.k8s.io/controller created
rolebinding.rbac.authorization.k8s.io/pod-lister created
clusterrolebinding.rbac.authorization.k8s.io/metallb-system:controller created
clusterrolebinding.rbac.authorization.k8s.io/metallb-system:speaker created
secret/webhook-server-cert created
service/webhook-service created
deployment.apps/controller created
daemonset.apps/speaker created
validatingwebhookconfiguration.admissionregistration.k8s.io/metallb-webhook-configuration created
```

※ 참고 : <https://metallb.universe.tf/installation/#installation-by-manifest>



# Check

remote > **kubectl get namespaces**

NAME	STATUS	AGE
default	Active	24d
ingress-nginx	Active	16h
kube-node-lease	Active	24d
kube-public	Active	24d
kube-system	Active	24d
metallb-system	Active	3m12s

remote > **kubectl get pods -o wide --namespace metallb-system**

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED	NODE	READINESS	GATES
controller-64cc46b9f9-vxqb2	1/1	Running	0	4m40s	10.233.103.83	worker2	<none>		<none>	
speaker-dwprs	1/1	Running	0	4m40s	192.168.100.201	worker1	<none>		<none>	
speaker-n8prv	1/1	Running	0	4m40s	192.168.100.202	worker2	<none>		<none>	
speaker-prhkx	1/1	Running	0	4m40s	192.168.100.200	master	<none>		<none>	

# Configuration

- L2 / BGP 유형에 따라 설정이 다르지만, 여기에서는 L2 기준으로 진행하겠다.

## metallb-ip.yaml

```
apiVersion: metallb.io/v1beta1
kind: IPAddressPool

metadata:
  name: lb-pool
  namespace: metallb-system

spec:
  addresses:
    - 192.168.100.240-192.168.100.250
```

## metallb-l2.yaml

```
apiVersion: metallb.io/v1beta1
kind: L2Advertisement

metadata:
  name: lb
  namespace: metallb-system

spec:
  ipAddressPools:
    - lb-pool
```

```
remote > cd kubernetes/05-Ingress-LoadBalancer/hands-on
```

```
remote > kubectl create -f metallb-ip.yaml
```

```
ipaddresspool.metallb.io/lb-pool created
```

```
remote > kubectl create -f metallb-l2.yaml
```

```
l2advertisement.metallb.io/lb created
```

※ 참고 : <https://metallb.universe.tf/configuration/#layer-2-configuration>



# Ready

- 앞에서 진행했던 ReplicaSet (Pods) 및 Service를 그대로 사용해서 LoadBalancer를 실습해보도록 하겠다.

```
remote > kubectl get replicaset -o wide
```

NAME	DESIRED	CURRENT	READY	AGE	CONTAINERS	IMAGES	SELECTOR
rs-web-v1	3	3	3	8m55s	node-web	whatwant/node-web:1.0	app=node-web-v1
rs-web-v2	3	3	3	8m52s	node-web	whatwant/node-web:2.0	app=node-web-v2

```
remote > kubectl get pods -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
rs-web-v1-ghbqs	1/1	Running	0	9m24s	10.233.110.29	worker1	<none>	<none>
rs-web-v1-m6gtz	1/1	Running	0	9m24s	10.233.103.79	worker2	<none>	<none>
rs-web-v1-rxpxz	1/1	Running	0	9m24s	10.233.103.80	worker2	<none>	<none>
rs-web-v2-cmvl2	1/1	Running	0	9m21s	10.233.103.82	worker2	<none>	<none>
rs-web-v2-k6src	1/1	Running	0	9m21s	10.233.110.30	worker1	<none>	<none>
rs-web-v2-qsh68	1/1	Running	0	9m21s	10.233.103.81	worker2	<none>	<none>

```
remote > kubectl get services -o wide
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE	SELECTOR
kubernetes	ClusterIP	10.233.0.1	<none>	443/TCP	24d	<none>
svc-node-web-v1	NodePort	10.233.24.241	<none>	80:30001/TCP	11s	app=node-web-v1
svc-node-web-v2	NodePort	10.233.26.196	<none>	80:30002/TCP	7s	app=node-web-v2

# Create LoadBalancer

- LoadBalancer 유형의 Service를 생성하는 것은 type 으로 명시만 해주면 된다.

## svc-node-web-lb.yaml

```
apiVersion: v1
kind: Service
metadata:
  name: svc-lb
spec:
  type: LoadBalancer

  ports:
  - name: http
    port: 80
    protocol: TCP
    targetPort: 8080

  selector:
    app: node-web-v1
```

```
remote > cd kubernetes/05-Ingress-LoadBalancer/hands-on
```

```
remote > kubectl create -f svc-node-web-lb.yaml
```

```
service/svc-lb created
```

```
remote > kubectl get services -o wide
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE	SELECTOR
kubernetes	ClusterIP	10.233.0.1	<none>	443/TCP	24d	<none>
svc-lb	LoadBalancer	10.233.46.241	192.168.100.240	80:31336/TCP	4s	app=node-web-v1
svc-node-web-v1	NodePort	10.233.24.241	<none>	80:30001/TCP	7h54m	app=node-web-v1
svc-node-web-v2	NodePort	10.233.26.196	<none>	80:30002/TCP	7h54m	app=node-web-v2

```
remote > curl -s http://192.168.100.240
```

```
You've hit rs-web-v1-m6gtz
```

# Describe LoadBalancer

```
remote > kubectl describe service svc-lb
```

```
Name: svc-lb
Namespace: default
Labels: <none>
Annotations: <none>
Selector: app=node-web-v1
Type: LoadBalancer
IP Family Policy: SingleStack
IP Families: IPv4
IP: 10.233.46.241
IPs: 10.233.46.241
LoadBalancer Ingress: 192.168.100.240
Port: http 80/TCP
TargetPort: 8080/TCP
NodePort: http 31336/TCP
Endpoints: 10.233.103.79:8080,10.233.103.80:8080,10.233.110.29:8080
Session Affinity: None
External Traffic Policy: Cluster
```

Events:

Type	Reason	Age	From	Message
----	-----	----	----	-----
Normal	IPAllocated	3m19s	metalb-controller	Assigned IP ["192.168.100.240"]
Normal	nodeAssigned	2m2s (x25 over 3m19s)	metalb-speaker	announcing from node "master" with protocol "layer2"




**sessionAffinity**  
**(Sticky Session)**



# Problem

Without Session Stickiness



With Session Stickiness



- 앞에서 실습한 LoadBalancer 를 조금 더 살펴보자.

```
remote > curl -s http://192.168.100.240
You've hit rs-web-v1-ghbqs

remote > curl -s http://192.168.100.240
You've hit rs-web-v1-rxpxz

remote > curl -s http://192.168.100.240
You've hit rs-web-v1-m6gtz

remote > curl -s http://192.168.100.240
You've hit rs-web-v1-rxpxz
```

- 접속할 때 마다 매번 바뀌는 Endpoints 를 볼 수 있다 → 만약 게시판과 같은 웹사이트라면 문제가 있을 수 있다!

※ 참고 : <https://www.imperva.com/learn/availability/sticky-session-persistence-and-cookies/>

# sessionAffinity (Sticky Session)

- 세션이 유지되는 동안 접속되는 Pod가 동일하도록 하기 위한 설정을 해보자.

## svc-node-web-lb-sa.yaml

```
apiVersion: v1
kind: Service
metadata:
  name: svc-lb-sa
spec:
  type: LoadBalancer

  ports:
    - name: http
      port: 80
      protocol: TCP
      targetPort: 8080

  selector:
    app: node-web-v1
```

```
sessionAffinity: ClientIP
sessionAffinityConfig:
  clientIP:
    timeoutSeconds: 3600
```

```
remote > cd kubernetes/05-Ingress-LoadBalancer/hands-on
```

```
remote > kubectl create -f svc-node-web-lb-sa.yaml
```

```
service/svc-lb-sa created
```

```
remote > kubectl get services -o wide
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE	SELECTOR
kubernetes	ClusterIP	10.233.0.1	<none>	443/TCP	24d	<none>
svc-lb	LoadBalancer	10.233.46.241	192.168.100.240	80:31336/TCP	49m	app=node-web-v1
svc-lb-sa	LoadBalancer	10.233.20.177	192.168.100.241	80:30566/TCP	25s	app=node-web-v1

```
remote > curl -s http://192.168.100.241
```

```
You've hit rs-web-v1-ghbqs
```

```
remote > curl -s http://192.168.100.241
```

```
You've hit rs-web-v1-ghbqs
```

```
remote > curl -s http://192.168.100.241
```

```
You've hit rs-web-v1-ghbqs
```

- `sessionAffinityConfig` 기본값은 10,800sec 이다. 즉, 지정하지 않아도 되는 설정이지만 공부를 위해서 추가해보았다.



**<https://kahoot.it/>**



# 자습 (복습)

▷ Ingress

▷ LoadBalancer