

Kubernetes - Beyond a Black Box

A humble peek into one level below your running application in production

Part II

About The Author

- **Hao (Harry) Zhang**
- **Currently:** Software Engineer, LinkedIn, Team Helix @ Data Infrastructure
- **Previously:** Member of Technical Staff, Applatix, Worked with Kubernetes, Docker and AWS
- **Previous blog series:** “Making Kubernetes Production Ready”
 - [Part 1](#), [Part 2](#), [Part 3](#)
 - Or just Google “[Kubernetes production](#)”, “[Kubernetes in production](#)”, or similar
- [Connect with me on LinkedIn](#)

Previously in Part I

- A high level idea about what is Kubernetes
- Components, functionalities, and design choice
 - API Server
 - Controller Manager
 - Scheduler
 - Kubelet
 - Kube-proxy
- Put it together, what happens when you do `kubectl create`
- [SlideShare link for Part I](#)

Outline - Part II

- An analysis of controlling framework design philosophy
 - **Choreography, not Orchestration**
 - **Level Driven, not Edge Triggered**
 - **Generalized Workload and Centralized Controller**
- Scheduling - Limitation and next steps
- Interfaces for production environments
- High level workload abstractions
 - Strength and limitations
- Conclusion

Part II

Controlling Framework

Choreography, Level-driven, Centralized Controller

Micro-service Choreography

- Orchestration: one gigantic controller trying to make everything correct at the same time
- Choreography: Desired state in cluster is achieved by **collaborations of separate autonomous entities reacting on changes of one or more API object(s) they are interested in**
 - Separation of concern
 - More flexibility to extend different types of semantics (CRD / TPR are great examples)
 - Develop a controller is easy!

More Readings:

- ❖ Third Party Resource (TPR): <https://kubernetes.io/docs/tasks/access-kubernetes-api/extend-api-third-party-resource/>
- ❖ Customer Resource Definition (CRD): <https://kubernetes.io/docs/tasks/access-kubernetes-api/extend-api-custom-resource-definitions/>
- ❖ TPR Design Spec: <https://github.com/kubernetes/community/blob/master/contributors/design-proposals/api-machinery/extending-api.md>

Micro-service Choreography

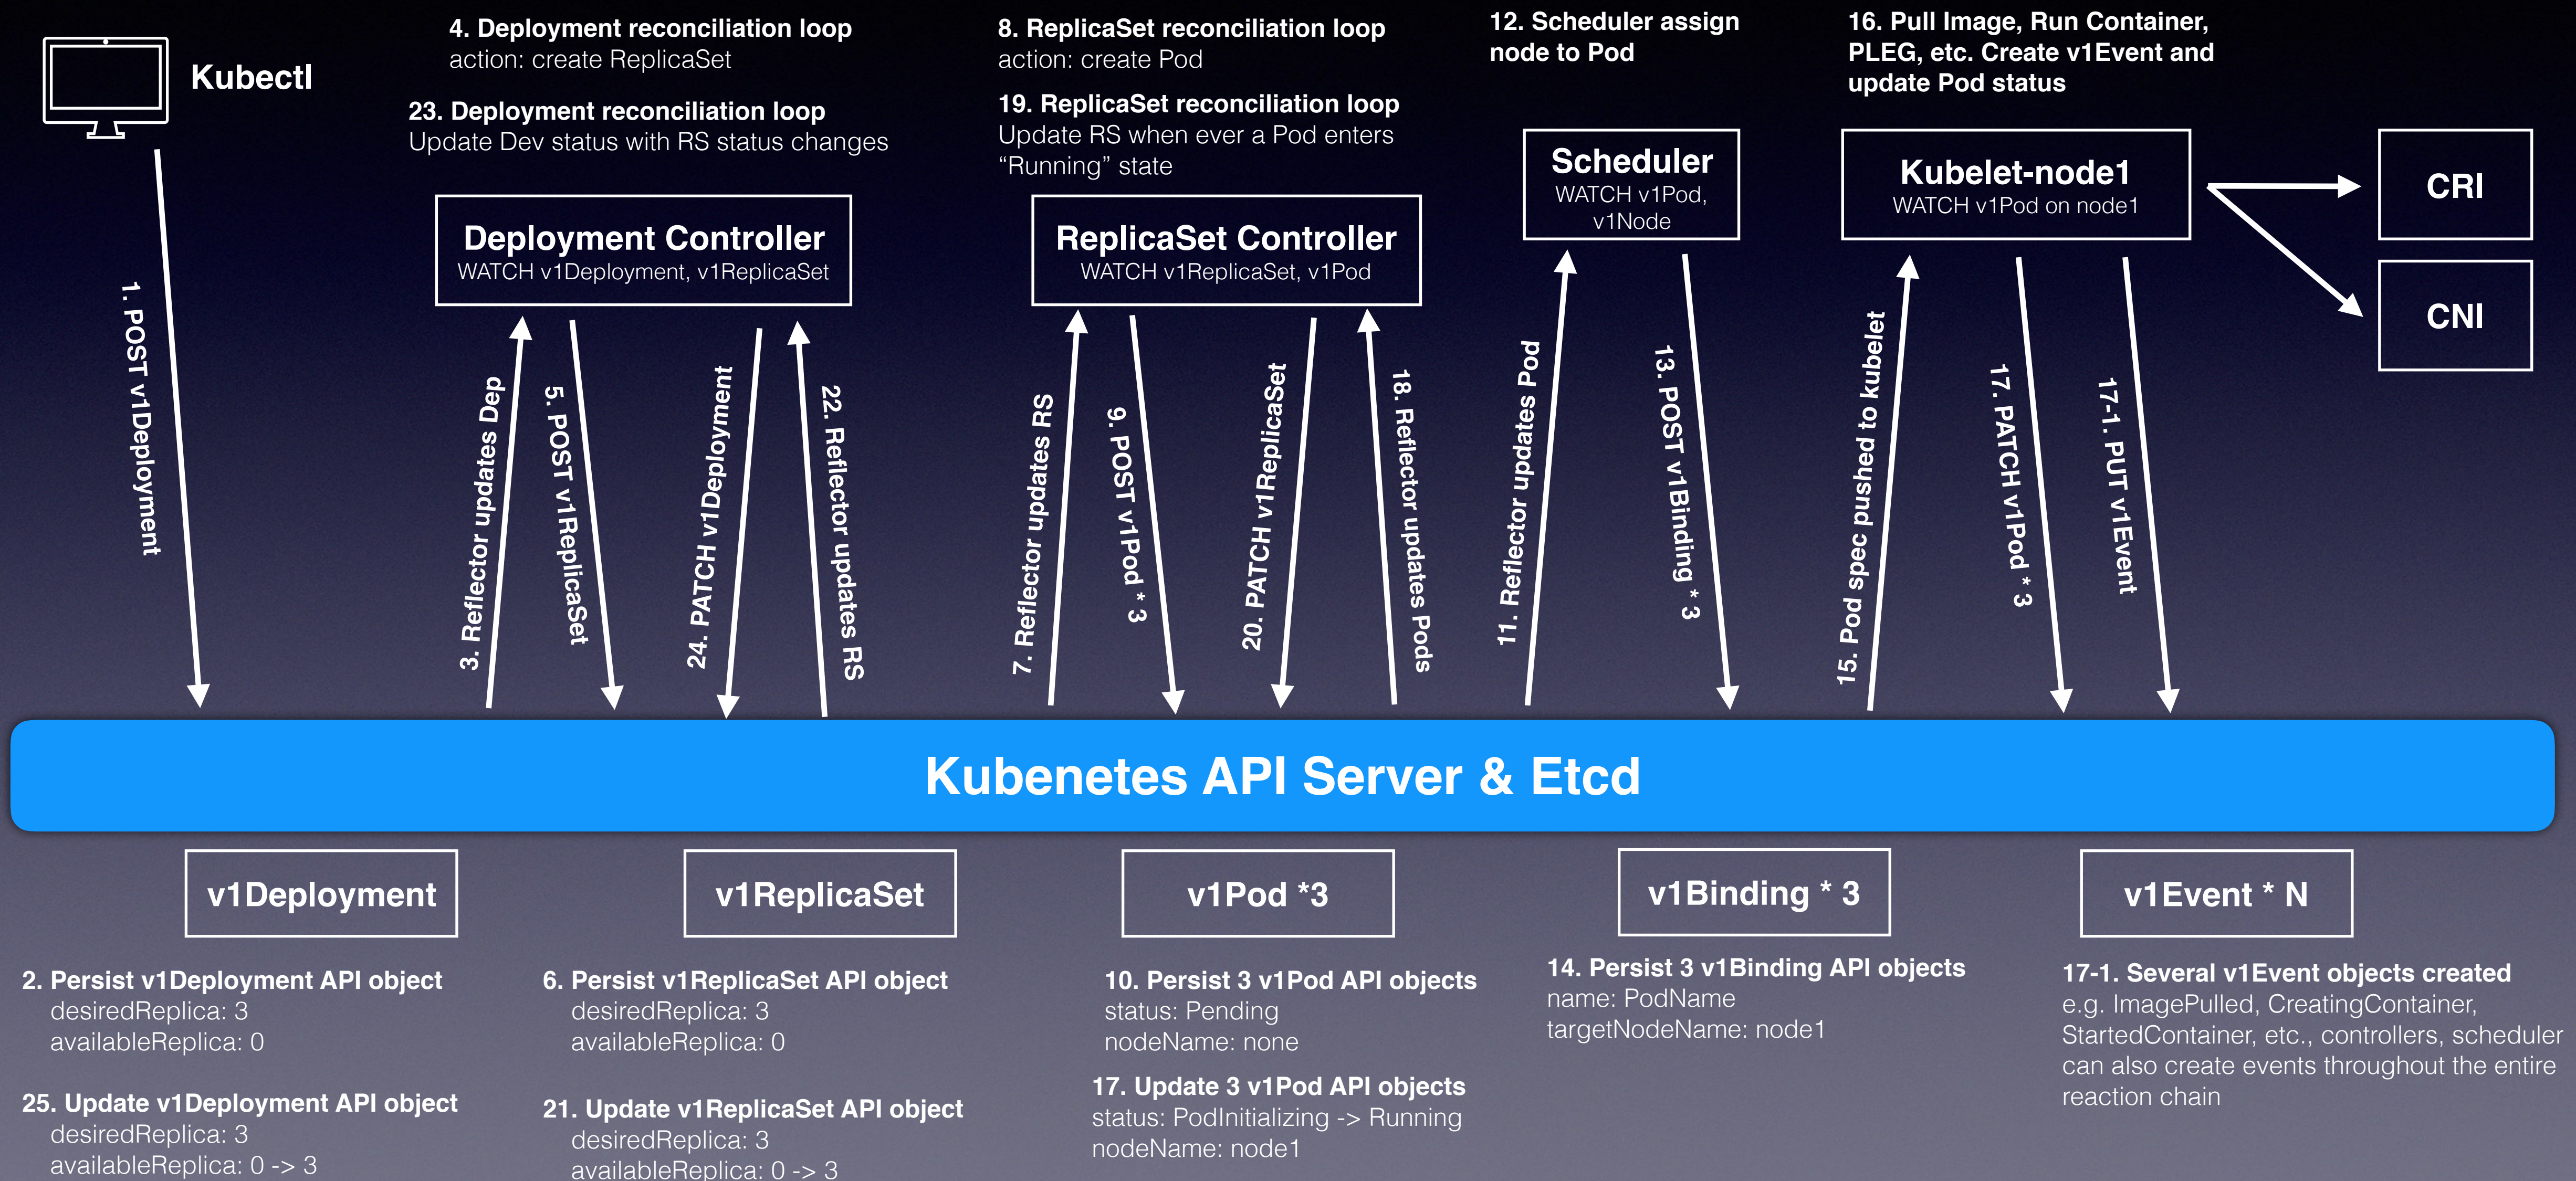
- Separation of Concern via **Micro-service Choreography** (Deployment example)
 - **ReplicaSet (RS)** ensures given # of Pod is always running
 - Create / Delete Pod API object if existing number is not matching spec
 - Scheduler knows where to put it; Kubelet knows Pod life cycle details
 - **Deployment (DEP)** provides higher level of semantics
 - **Scale:** tell RS to change # of Pods by modifying RS spec, RS knows what to do
 - **Upgrade:** delete current RS, create new RS with new Pod spec
 - **Rolling-Upgrade:** create new RS with new Pod spec, scale down old RS by 1, scale out new RS by 1, repeat it until no old Pod remains (Same for roll back)
 - **HorizontalPodAutoscaler (HPA)** manipulates deployment
 - It polls Pod metrics (from a source such as Heapster or Prometheus), compare it with user defined metric specs, and make scaling decision
 - Tell DEP to scale up/down by modifying DEP spec, DEP controller knows what to do

Micro-service Choreography

- See Next Slide: An illustration about how different autonomous entities reacting on API objects they are interested in can move cluster to a desired state
 - Using Deployment as example, duplicated from Part I

Micro-service Choreography

Note: this is a rather high-level idea about Kubernetes reaction chain. It hides some details for illustration purpose and is different than actual behavior



Micro-service Choreography

- Choreography - **some down sides**
 - Long reaction chain
 - Multiple layers of reconciliation stages can be error prone (i.e. racing conditions, component down, etc), it has been improved a lot in late Kubernetes versions
 - Latency could be long as pipeline needs to persist after every stage, might not fit high QoS requirements
 - Debugging is hard
 - i.e. I forget to update image in registry
 - GET Deployment — found desired replica is not up
 - LIST Pods with label — found image pull backoff in container status
 - LIST events with involved object as this Pod — found message “Image does not exist”
 - Operations can be heavy if you want to automate failure recovery / detailed error reporting, unless you cache a lot (Efficient caching is another huge topic...)

Level-Driven Control

- State is considered as “**level**” while state change resembles “**edge**”
- Desired state and current state are persisted
- Controller can always re-check object state and make action to move object towards desired state
- Atomic ListAndWatch (see Part I) assures controllers react based on state changes pushed from API server (watch) while not miss any event (list)

More Readings:

♣ Kubernetes: edge vs level triggered logics: <https://speakerdeck.com/thockin/edge-vs-level-triggered-logic>

Centralized Controller

- **Generalized Workload and Centralized Controller**
 - Kubernetes generalize applications into limited types of semantics
 - Job, Deployment, Node, StatefulSet, DaemonSet, ...
 - Reduced control overhead, i.e., 1 type of controller manages all instances of workloads in one category
 - i.e. Deployment controller will control all deployments in the cluster
 - Compared with 1 controller controls 1 deployment instance ($O(n)$ space complexity), Kubernetes' control overhead is constant
 - Reconciliation control loops: things will ultimately become correct
 - **Some down sides:** Stateful applications are hard to generalize (More discussions later)

Centralized Controller

- Two opposite design choices from state-of-art cluster management frameworks
 - Apache Hadoop Yarn, Apache Mesos
 - User have full freedom to implement their workload controlling logics
 - One app controller per app, all controllers talk to master
 - Kubernetes
 - Pre-defined very generic workload abstractions
 - Workloads of same type share 1 controller
 - Centralized management of controllers with optimized master access

Scheduling

Limitations and Next Steps

Scheduling

- Default sequential scheduler - **Limitation**
 - Scoring system can be hard to tune
 - Sequential scheduling might not be ideal for batch workloads
 - Assumes launching more is ALWAYS better than launching less
 - I need 3 replicas to form a quorum, then what's the point of starting 2 upon insufficient resource?
- No global re-balancing mechanism
 - Upon new node join
 - Moving things around can reduce node resource pressure as we over-provision
 - But this rebalance would make scale-down even harder
 - Upon insufficient resource (moving things around might fit)
 - Hacks available such as re-scheduling and Kubelet preemption

Scheduling

- Some advanced scheduling problems people are trying to solve:
 - If m Pods cannot fit onto n Nodes, how to choose which ones to run
 - If not all Pods in an application can get scheduled
 - Can we just schedule some of them?
 - What if an application needs to create other resources?
 - i.e. workflow has a big parallel step which cannot be serialized, if this step cannot fit in, it'd be better to fail the entire workflow
 - Cluster-wide resource isolation mechanism other than ResourceQuota per namespace
 - Bin-pack apps and scale back to release some resources

More Readings:

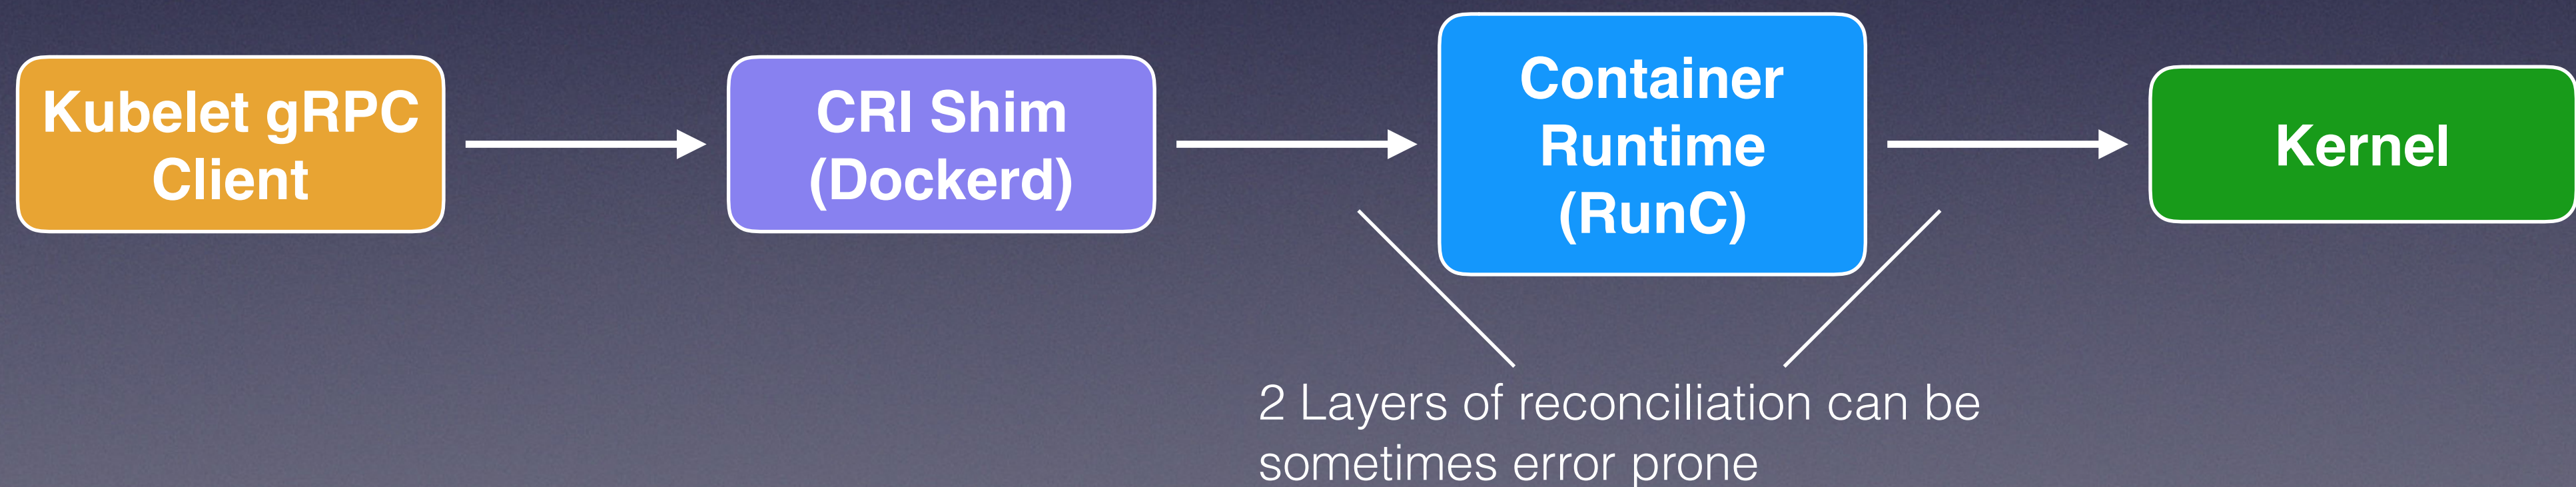
- ❖ Resource Sharing / scheduling design spec: https://docs.google.com/document/d/1-H2hnZap7gQivcSU-9j4ZrJ8wE_WwcfOkTeAGjzUyLA
- ❖ Scheduling Feature Proposals: <https://github.com/kubernetes/community/tree/master/contributors/design-proposals/scheduling>

Environment Interfaces

CNI, CRI and Flex Volume

Production Interfaces

- Container Runtime Interfaces (CRI)
 - Implemented based on Open Container Initiative (OCI)
 - Runtime Service for container lifecycle management
 - Image Service for image management
 - Runtime plugin can be customized
 - You can even write your plug-in for VM



More Readings:

✿ Docker reconciliation bug: <https://github.com/moby/moby/issues/32413>

Production Interfaces

- Container Network Interface (CNI)
 - Pluggable interface for cluster networking layer
 - Used to setup/teardown Pod network
- Flex Volume (user-defined volumes)
 - Interface for data volume plugin
 - Need to implement methods such as attach/detach, mount/unmount
- With these CRI, CNI, and Flex Volume, you can make kubelet a “dumb”, environment agnostic worker, which is extremely flexible to fit any production environment

Production Interfaces

More Readings:

❖ CRI

- ❖ Official blog introducing CRI: <http://blog.kubernetes.io/2016/12/container-runtime-interface-cri-in-kubernetes.html>
- ❖ CRI Spec: <https://github.com/kubernetes/community/blob/master/contributors/devel/container-runtime-interface.md>
- ❖ CRI Container Stats Proposal: <https://github.com/kubernetes/community/blob/master/contributors/devel/cri-container-stats.md>
- ❖ Open Container Initiative (OCI): <https://www.opencontainers.org>
- ❖ Open Container Initiative (OCI) Runtime/Image spec, and RunC: <https://github.com/opencontainers>

❖ CNI

- ❖ Pod Networking Design Proposal: <https://github.com/kubernetes/community/blob/master/contributors/design-proposals/network/networking.md>
- ❖ Kubernetes Networking Plugin Introduction: <https://kubernetes.io/docs/concepts/cluster-administration/network-plugins/#cni>
- ❖ Linux CNI Spec: <https://github.com/containernetworking/cni/blob/master/SPEC.md#network-configuration>
- ❖ CNCF CNI Project: <https://github.com/containernetworking/cni>
- ❖ Kubelet CNI Usage: <https://github.com/kubernetes/kubernetes/blob/master/pkg/kubelet/network/cni/cni.go>
- ❖ Kubenet - Kubernetes' default network plugin: <https://github.com/kubernetes/kubernetes/tree/master/pkg/kubelet/network/kubenet>

❖ Volume

- ❖ Flex Volume: <https://github.com/kubernetes/community/blob/master/contributors/devel/flexvolume.md>
- ❖ Kubernetes Volume Documentation: <https://kubernetes.io/docs/concepts/storage/volumes/>

Workload Abstractions

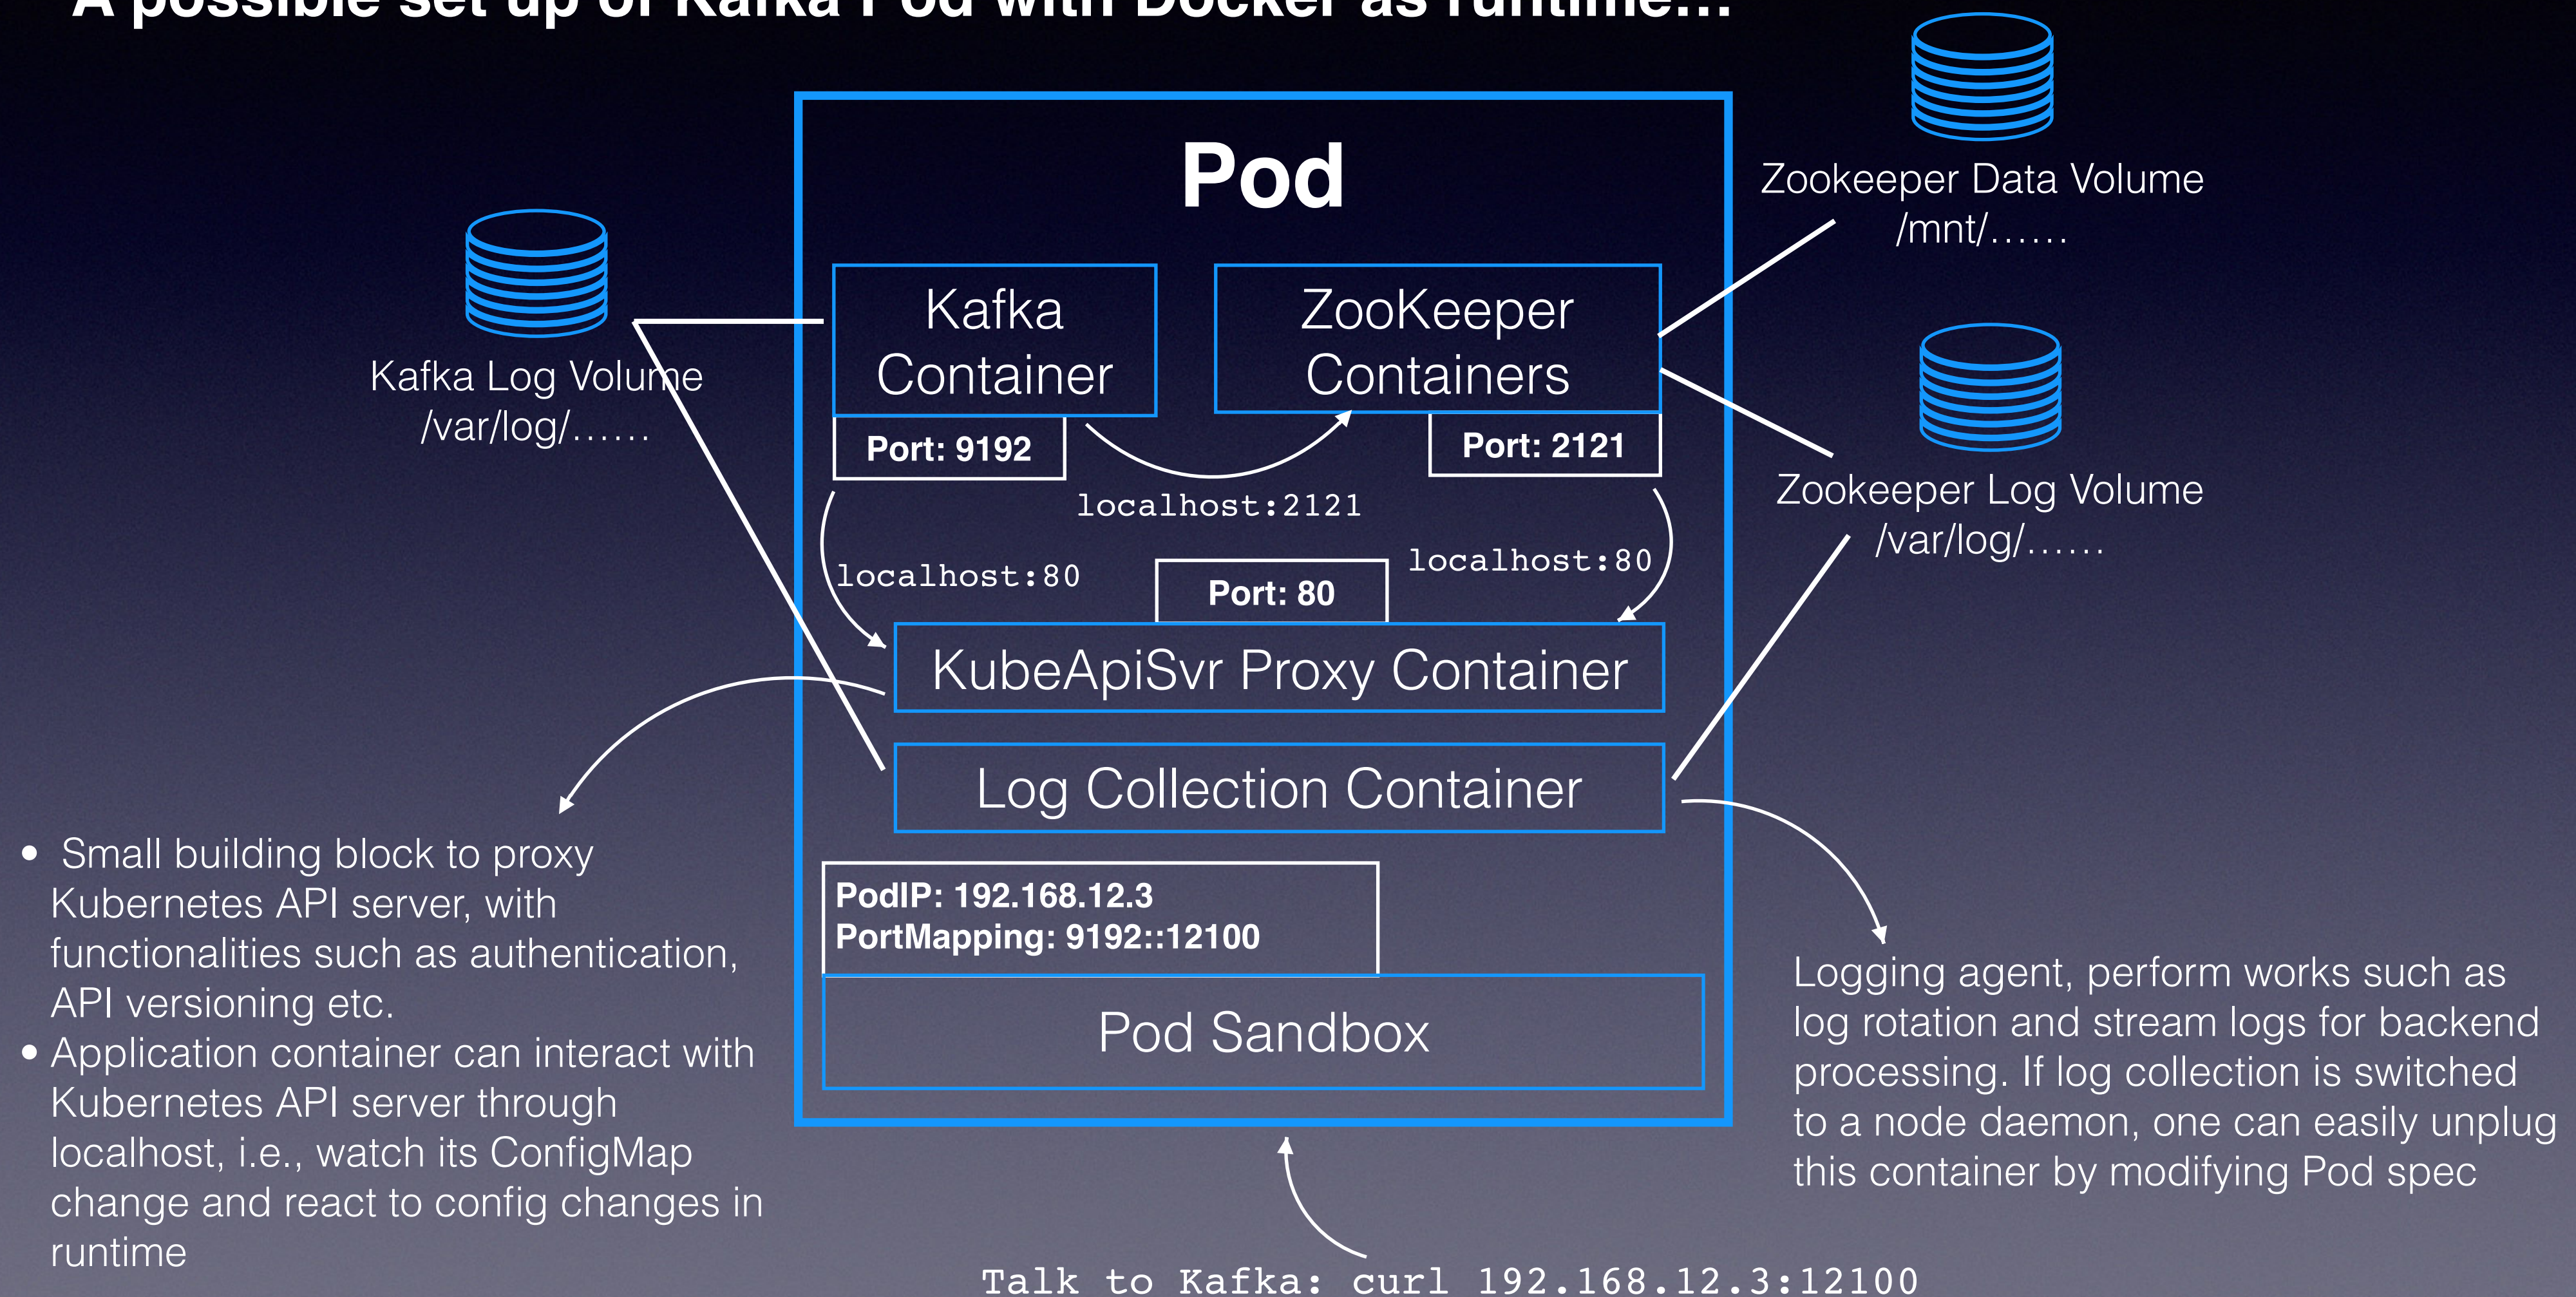
Pod, Job, Deployment, DaemonSet, StatefulSet - Strength
and limitations

Workload Abstractions

- Pod - Basic Workload Unit in Kubernetes
 - An isolated environment with resource constraints
 - With docker as run time, it is a group of containers under Infra container Cgroup
 - User can implement their own runtime through CRI
 - Pod = atomic scheduling unit, ensures co-location
 - Container = single-function building block running in isolated env
 - Just modify Pod spec to plug/unplug functionalities, no code change
 - Spin up in milliseconds (it's just a process)
 - Containers in Pod can share directory, volume, localhost, etc.
 - Crashed container will be restarted on same node

Workload Abstractions

A possible set up of Kafka Pod with Docker as runtime...



Workload Abstractions

- Job
 - Run to complete (i.e. container exit code is 0)
 - Cron job, batch job, etc...
- Deployment
 - Maintain replicas of **stateless** applications (not managing volume)
 - High-level interfaces such as rollout, rollback
- DaemonSet
 - One replica per node, primary use cases include:
 - Log collection daemon (fluentd)
 - Node monitoring daemon (node-exporter)
 - Cluster storage daemon (gclusterd)

Workload Abstractions

- StatefulSet
 - Use case prototypes:
 - Quorum with leader election: MongoDB, Zookeeper, Etcd
 - De-centralized quorum: Cassandra
 - Active-active (multiple masters): Galera
 - Besides features of deployment, StatefulSet also provides:
 - Every replica has persistent identifier for network (Pod name is formatted as “podName- $\{0..N-1\}$ ”), which might help identifying peers
 - Every replica has persistent storage (data persist even after deletion)
 - Supports automatic storage provisioning
 - Ordered deploy, shutdown and upgrade (from $\{0..N-1\}$)

Workload Abstractions

More Readings:

- ✿ Typically used workloads
 - ✿ Pod: <https://kubernetes.io/docs/concepts/workloads/pods/pod/>
 - ✿ Deployment: <https://kubernetes.io/docs/concepts/workloads/controllers/deployment/>
 - ✿ Job: <https://kubernetes.io/docs/concepts/workloads/controllers/jobs-run-to-completion/>
 - ✿ CronJob: <https://kubernetes.io/docs/concepts/workloads/controllers/cron-jobs/>
 - ✿ Daemonset: <https://kubernetes.io/docs/concepts/workloads/controllers/daemonset/>
- ✿ Other Kubernetes Concepts (Network/Config/Volume, etc.): <https://kubernetes.io/docs/concepts/>
- ✿ Design Specs
 - ✿ Workloads Design Specs (Interfaces, behaviors and updates): <https://github.com/kubernetes/community/tree/master/contributors/design-proposals/apps>
 - ✿ Autoscaling Design Specs: <https://github.com/kubernetes/community/tree/master/contributors/design-proposals/autoscaling>
- ✿ Very good example of running master-slave MySQL example on Kubernetes using StatefulSet: <https://kubernetes.io/docs/tasks/run-application/run-replicated-stateful-application/>

Abstraction Limitations

- Some thoughts about Kubernetes workload abstractions
 - **Pod** and **Job** are perfect for small execution unit
 - **Deployment** can fit most use cases of stateless application
 - **DaemonSet** can fit most use cases of one-per-node applications
 - **StatefulSet** might be useful for small-medium scale peer-to-peer apps
 - i.e. Leader election can be done through DNS, and does not need StatefulSet controller's help
 - Over-generalized StatefulSet controller is application-agnostic and therefore cannot control complicated application state
- Good thing is that Kubernetes makes it easy enough to define customer resources and write controllers

Abstraction Limitations

- Complicated Workload Case 1: **Workflow**
 - Stateful and run-to-complete
 - There used to be an official DAG workflow implementation, but was finally moved out from core API
 - Even usually defined as DAGs, Workflow can be complicated in many different ways and is very hard to generalize as core API
 - i.e. A deep-learning workflow can be totally different from a DevOps workflow
 - It's not always just DAG, it can also contain FOR-loops, If-Else, etc
 - Impl Discussions: <https://github.com/kubernetes/kubernetes/pull/17787>
 - Design Spec: <https://github.com/kubernetes/kubernetes/pull/18827>
 - Implementation: <https://github.com/kubernetes/kubernetes/pull/24781>
 - Discussion to Remove: <https://github.com/kubernetes/kubernetes/issues/25067>

Abstraction Limitations

- Complicated Workload Case 2: **Semi-stateful Apps**
 - DevOps use case: a pool of workers working on building docker images
 - Need a persistent volume as graph storage to cache docker image layers (if layer is not updated, don't need to pull from remote)
 - Data (image layers) can be huge so not a good idea to use node's root storage or memory
 - Losing data is fine as this is just a cache
 - But you can also say, it's persisting data so its stateful...
 - StatefulSet is too heavy but ReplicationSet does not support dynamically provisioned data volumes

Abstraction Limitations

- Complicated Workload Case 3: **Sharded DB with Master/Slave of Each Shard**
 - Master RW, Slave RO
 - Master usually handle more workload
 - Need global view to balance load among multiple such applications
 - Node will suffer if too many replica on it become master
 - StatefulSet is NOT application-aware so additional work is needed
 - When a particular shard has request spike, what's better?
 - Possible action 1: Scale **horizontally** and re-shard
 - Horizontal-scaler might help
 - Overhead in data migration and re-sharding
 - Scale back is also challenging
 - Possible action 2: Scale **vertically** (up to node capacity) and evict
 - Might remove less important Pods from node and they can hopefully get re-scheduled

Conclusion

What makes Kubernetes successful

Take-aways from Kube's Design

- API Server and Versioned API Groups
 - Easy upgrade / backward compatibility
 - API server performs conversion between client and metadata store
- Protect cluster, every where
 - Throttle, resource quota limit @ API server
 - Important production question: **Given a pre-provisioned metadata store, what is a reasonable amount of mutating / non-mutating operations I can smoothly handle in parallel based on my SLA?**
 - QPS, exponential backoff with jittered retry @ client

Take-aways from Kube's Design

- Optimize Etcd READs
 - Cache @ API server - serve all READs from memory
 - It's just meta data, size can be reasonable for memory
 - Shared Informer - aggregate READs from all controllers
- Micro-service choreography
 - Everyone focus on their own API objects and different pipelines are formed automatically
 - Atomic commits on single API objects, reconciliation loops will finally make things right
- Level-triggered control
 - Controller always go back and assert state, so nothing can be missing

Why Cluster Management

- As tedious as endless on-calls
- As sexy as and as important as an orchestrator that
 - Increases resource utilization and save money
 - Automates resource provisioning / scheduling / scaling (both up and down) / failure detection and recovery
 - Provides methods to debug from outside the cluster
 - i.e. execute debug commands in your container
 - Plugs solutions such as Logging, Monitoring, Dashboard, Security, Admission Control, etc.
 - Release labor for feature development

Why is Kubernetes Successful

- Great abstraction for plug-and-play simple workload
 - Writing control logics can somewhat be a burden, especially for startups, and Kubernetes made it much easier
- Environment-agnostic interfaces
 - CRI, CNI, Flex Volume makes it possible to handle hybrid environment
- Native support for popular public cloud
 - Out-of-box resource provisioning and management
- Plug-and-play cluster management solutions from community
 - Cluster autoscaler, logging, monitoring, admission, etc.
 - Just run `kubectl create` you will have your app

Borg, Omega, Kubernetes

- Borg @ Google: <https://research.google.com/pubs/pub43438.html>
- Omega @ Google: <https://research.google.com/pubs/pub41684.html>
- From Borg, Omega to Kubernetes: <http://queue.acm.org/detail.cfm?id=2898444>