

**5<sup>th</sup>**  
**Week**

# 다섯 번째 뵙겠습니다 ?!

▷ 잠시만 기다렸다가 30분 되면 시작하겠습니다~^^

▷ 지난 주에는 조금 일찍 끝났으니, 이번 주에는 ... ?!

- 신나게 달려봅시다~~~!!!

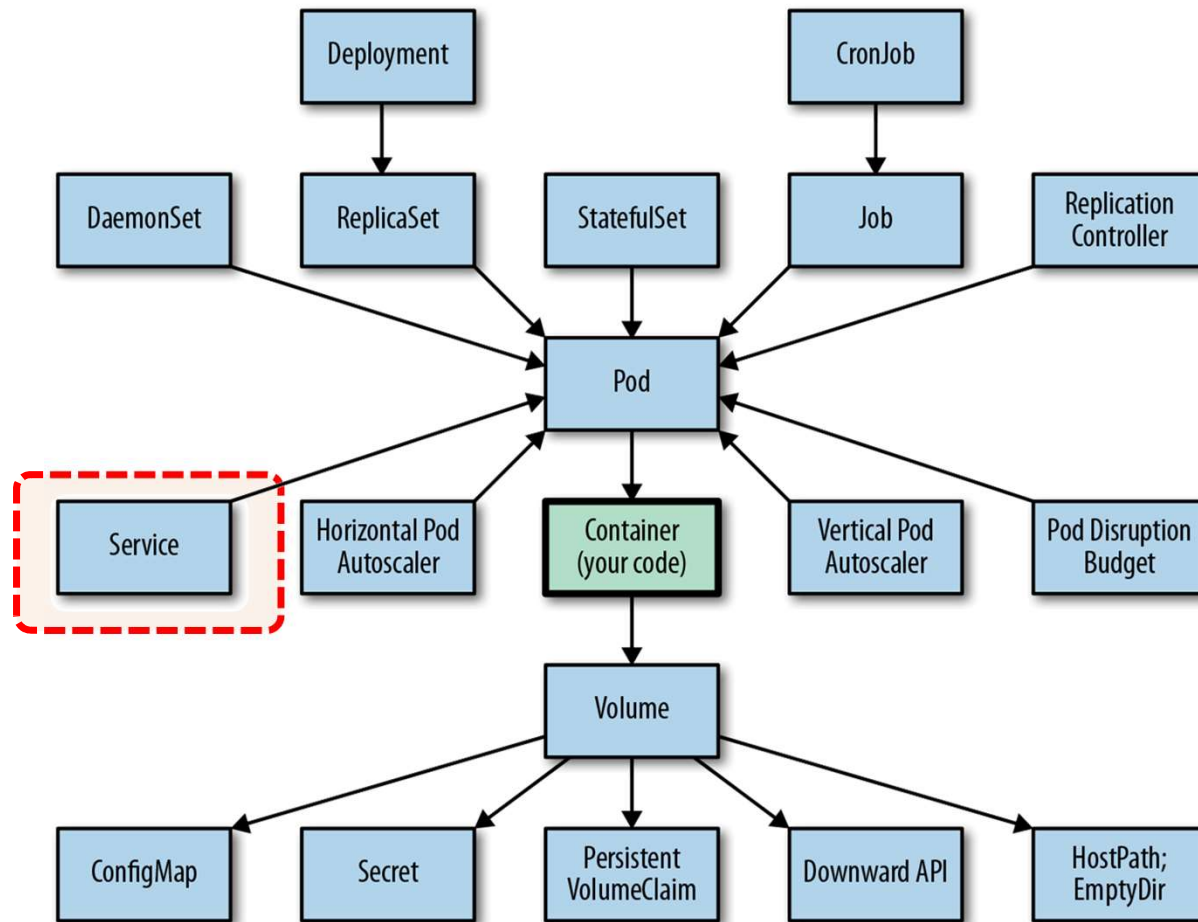
▷ Camera는 가급적 켜 주시면 대단히 감사하겠습니다 !!!

- 너무 부끄러우면 Snap Camera를 사용하시는 것까지는~ ^^

▷ 오늘 수업 자료는 아래 링크에서 다운로드 받으실 수 있어요.

- <https://github.com/whatwant-school/kubernetes>

## Service



※ 참고 : <https://www.oreilly.com/library/view/kubernetes-patterns/9781492050278/ch01.html>



# 지난 수업 복습

**<https://kahoot.it/>**



# **Service**

**Ingress / LoadBalancer**





# **Flip Learning**

**(Service - Ingress / LoadBalancer)**

**권준혁 님**



# **Kubernetes**

## **Network - Ingress**

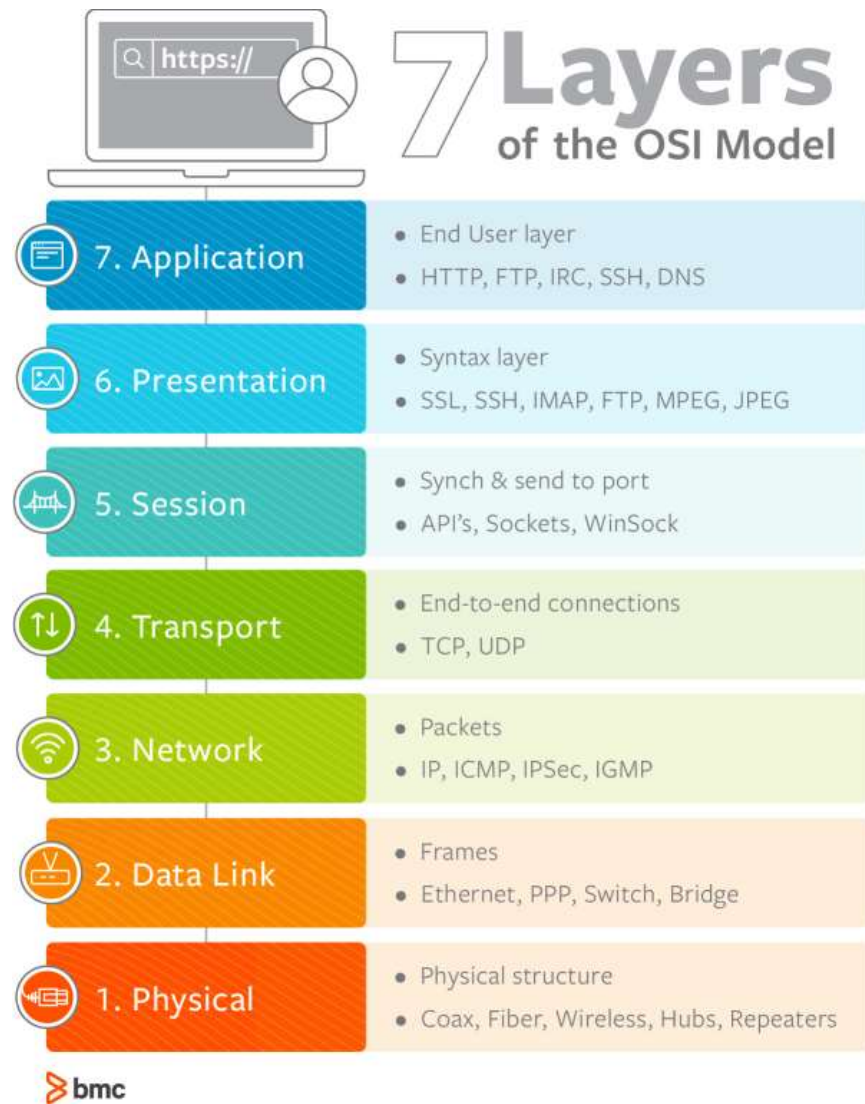
# Ingress Controller

- Ingress는 cluster 외부에서 내부 Service로 HTTP와 HTTPS 경로 노출
- 클라이언트가 요청한 호스트와 경로에 따라 요청을 전달할 서비스가 결정
  - . HTTP(S)기반의 L7 LoadBalancing 기능을 제공하는 컴포넌트
- Ingress 구현체 : 각 구현체마다 설정 방법 차이 존재
  - . 구글 클라우드 : <https://github.com/kubernetes/ingress-gce>
  - . 오픈소스 (Ingress NGINX Controller) : <https://github.com/kubernetes/ingress-nginx>
  - . 오픈소스 (NGINX Ingress Controller) : <https://github.com/nginxinc/kubernetes-ingress>
  - . 상용 (F5 BIG IP Controller) : <http://clouddocs.f5.com/products/connectors/k8s-bigip-ctrlr>



※ 참고 : <https://kubernetes.io/ko/docs/concepts/services-networking/ingress/>

# OSI 7 Layers



※ 참고 : <https://www.bmc.com/blogs/osi-model-7-layers/>

# Ingress Controller Install

Kubernetes에서 배포하고 있는 패키지로 설치해보는데, 아래 링크에서 bare-metal 설치 방법을 찾아보면 된다

- <https://github.com/kubernetes/ingress-nginx/blob/main/docs/deploy/index.md#bare-metal-clusters>

```
remote > kubectl apply -f https://raw.githubusercontent.com/kubernetes/ingress-nginx/controller-v1.3.0/deploy/static/provider/baremetal/deploy.yaml

namespace/ingress-nginx created
serviceaccount/ingress-nginx created
serviceaccount/ingress-nginx-admission created
role.rbac.authorization.k8s.io/ingress-nginx created
role.rbac.authorization.k8s.io/ingress-nginx-admission created
clusterrole.rbac.authorization.k8s.io/ingress-nginx created
clusterrole.rbac.authorization.k8s.io/ingress-nginx-admission created
rolebinding.rbac.authorization.k8s.io/ingress-nginx created
rolebinding.rbac.authorization.k8s.io/ingress-nginx-admission created
clusterrolebinding.rbac.authorization.k8s.io/ingress-nginx created
clusterrolebinding.rbac.authorization.k8s.io/ingress-nginx-admission created
configmap/ingress-nginx-controller created
service/ingress-nginx-controller created
service/ingress-nginx-controller-admission created
deployment.apps/ingress-nginx-controller created
job.batch/ingress-nginx-admission-create created
job.batch/ingress-nginx-admission-patch created
ingressclass.networking.k8s.io/nginx created
validatingwebhookconfiguration.admissionregistration.k8s.io/ingress-nginx-admission created
```

# Ingress Controller Check

```
remote > kubectl get namespaces
```

| NAME          | STATUS | AGE   |
|---------------|--------|-------|
| default       | Active | 23d   |
| ingress-nginx | Active | 2m19s |
| ...           |        |       |

```
remote > kubectl get pods -o wide --namespace kube-system
```

| NAME                | READY | STATUS  | RESTARTS       | AGE | IP              | NODE    | NOMINATED NODE | READINESS GATES |
|---------------------|-------|---------|----------------|-----|-----------------|---------|----------------|-----------------|
| ...                 |       |         |                |     |                 |         |                |                 |
| nginx-proxy-worker1 | 1/1   | Running | 11 (5m42s ago) | 23d | 192.168.100.201 | worker1 | <none>         | <none>          |
| nginx-proxy-worker2 | 1/1   | Running | 9 (5m34s ago)  | 23d | 192.168.100.202 | worker2 | <none>         | <none>          |
| ...                 |       |         |                |     |                 |         |                |                 |

```
remote > kubectl get pods -o wide --namespace ingress-nginx
```

| NAME                                      | READY | STATUS    | RESTARTS | AGE   | IP            | NODE    | NOMINATED NODE | READINESS GATES |
|---|-------|-----------|----------|-------|---------------|---------|----------------|-----------------|
| ingress-nginx-admission-create-6j9cv      | 0/1   | Completed | 0        | 3m49s | 10.233.103.75 | worker2 | <none>         | <none>          |
| ingress-nginx-admission-patch-qmwvf       | 0/1   | Completed | 1        | 3m49s | 10.233.110.24 | worker1 | <none>         | <none>          |
| ingress-nginx-controller-77cb5dbf4d-k5qmr | 1/1   | Running   | 0        | 3m49s | 10.233.110.25 | worker1 | <none>         | <none>          |

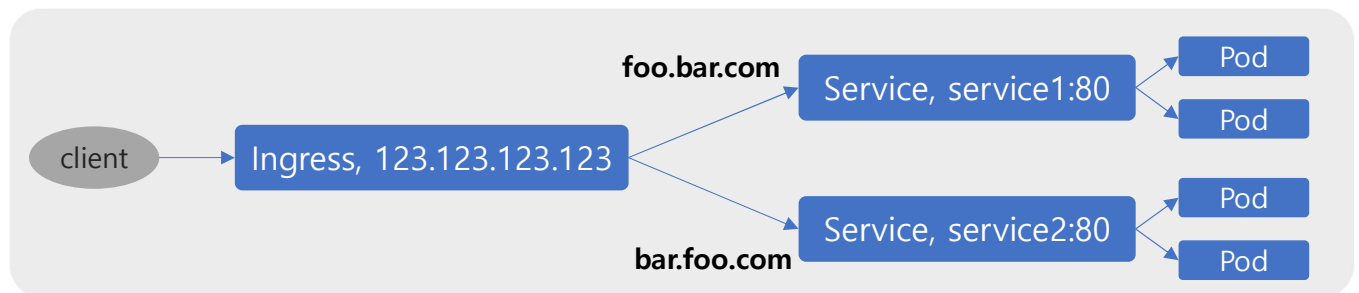
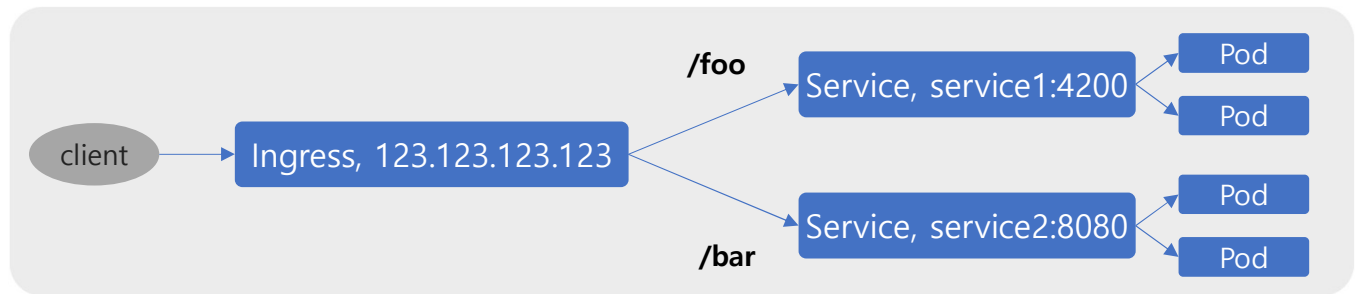

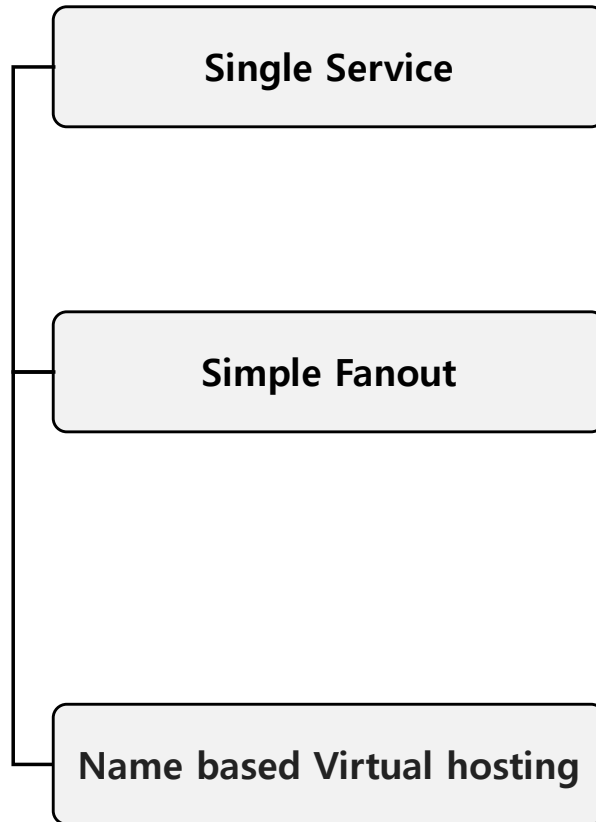
```
remote > kubectl get serviceaccounts -o wide --namespace ingress-nginx
```

| NAME                    | SECRETS | AGE   |
|-------------------------|---------|-------|
| default                 | 1       | 4m28s |
| ingress-nginx           | 1       | 4m28s |
| ingress-nginx-admission | 1       | 4m28s |





# Types of Ingress



※ 참고 : <https://kubernetes.io/docs/concepts/services-networking/ingress/>



# Ready: container images

별개의 사이트 구분을 할 수 있도록 웹서비스를 2종 준비했다.

기존 `node-web:1.0`과 구분이 될 정도만 살짝 변경한 후 `node-web:2.0` 버전으로 DockerHub에 push 했다.

```
const http = require('http');
const os = require('os');

console.log("node-web server starting...");

var handler = function(request, response) {
  console.log("Received request from " + request.connection.remoteAddress);
  response.writeHead(200);
  response.end("You've hit " + os.hostname() + "(Ver2.0)\n");
};

var www = http.createServer(handler);
www.listen(8080);
```

**app.js**

```
FROM node:latest
ADD app.js /app.js
ENTRYPOINT ["node", "app.js"]
```

**Dockerfile**

DockerHub에 이미 있는 만들어 놓은  
이미지를 사용해도 좋다.

```
remote > git clone https://github.com/whatwant-school/kubernetes.git
remote > cd kubernetes/05-Ingress-LoadBalancer/hands-on

remote > docker build -t node-web:2.0 .

remote > docker tag node-web:2.0 <user-id>/node-web:2.0

remote > docker push <user-id>/node-web:2.0
```

# Ready: ReplicaSet YAML

웹사이트 2종을 각각 띄울 `ReplicaSet`을 준비하자.

```
apiVersion: apps/v1
kind: ReplicaSet

metadata:
  name: rs-web-v1

spec:
  replicas: 3
  selector:
    matchLabels:
      app: node-web-v1

  template:
    metadata:
      labels:
        app: node-web-v1

    spec:
      containers:
        - name: node-web
          image: whatwant/node-web:1.0
          ports:
            - containerPort: 8080
          imagePullPolicy: Always
```

**rs-web-v1.yaml**

```
apiVersion: apps/v1
kind: ReplicaSet

metadata:
  name: rs-web-v2

spec:
  replicas: 3
  selector:
    matchLabels:
      app: node-web-v2

  template:
    metadata:
      labels:
        app: node-web-v2

    spec:
      containers:
        - name: node-web
          image: whatwant/node-web:2.0
          ports:
            - containerPort: 8080
          imagePullPolicy: Always
```

**rs-web-v2.yaml**

# Ready: Create ReplicaSet

앞에서 작성해 놓은 ReplicaSet을 이용해서 Pods를 생성하자.

```
remote > git clone https://github.com/whatwant-school/kubernetes.git
remote > cd kubernetes/05-Ingress-LoadBalancer/hands-on
```

```
remote > kubectl create -f rs-web-v1.yaml
```

```
replicaset.apps/rs-web-v1 created
```

```
remote > kubectl create -f rs-web-v2.yaml
```

```
replicaset.apps/rs-web-v2 created
```

```
remote > kubectl get pods -o wide
```

| NAME            | READY | STATUS  | RESTARTS | AGE   | IP            | NODE    | NOMINATED NODE | READINESS GATES |
|-----------------|-------|---------|----------|-------|---------------|---------|----------------|-----------------|
| rs-web-v1-ghbqs | 1/1   | Running | 0        | 2m25s | 10.233.110.29 | worker1 | <none>         | <none>          |
| rs-web-v1-m6gtz | 1/1   | Running | 0        | 2m25s | 10.233.103.79 | worker2 | <none>         | <none>          |
| rs-web-v1-rxpxz | 1/1   | Running | 0        | 2m25s | 10.233.103.80 | worker2 | <none>         | <none>          |
| rs-web-v2-cmvl2 | 1/1   | Running | 0        | 2m22s | 10.233.103.82 | worker2 | <none>         | <none>          |
| rs-web-v2-k6src | 1/1   | Running | 0        | 2m22s | 10.233.110.30 | worker1 | <none>         | <none>          |
| rs-web-v2-qsh68 | 1/1   | Running | 0        | 2m22s | 10.233.103.81 | worker2 | <none>         | <none>          |

# Ready: Create Service

`ClusterIP` 유형으로 해도 되지만, 여기에서는 `NodePort` 유형으로 `Service`를 만들어보자.

```
apiVersion: v1                                svc-node-web-nodeport-v1.yaml
kind: Service

metadata:
  name: svc-node-web-v1

spec:
  type: NodePort

  ports:
  - port: 80
    targetPort: 8080
    nodePort: 30001

  selector:
    app: node-web-v1
```

```
apiVersion: v1                                svc-node-web-nodeport-v2.yaml
kind: Service

metadata:
  name: svc-node-web-v2

spec:
  type: NodePort

  ports:
  - port: 80
    targetPort: 8080
    nodePort: 30002

  selector:
    app: node-web-v2
```

```
remote > cd kubernetes/05-Ingress-LoadBalancer/hands-on
```

```
remote > kubectl create -f svc-node-web-nodeport-v1.yaml
```

```
remote > kubectl create -f svc-node-web-nodeport-v2.yaml
```

```
remote > kubectl get services -o wide
```

| NAME            | TYPE      | CLUSTER-IP    | EXTERNAL-IP | PORT(S)      | AGE | SELECTOR        |
|-----------------|-----------|---------------|-------------|--------------|-----|-----------------|
| kubernetes      | ClusterIP | 10.233.0.1    | <none>      | 443/TCP      | 24d | <none>          |
| svc-node-web-v1 | NodePort  | 10.233.24.241 | <none>      | 80:30001/TCP | 11s | app=node-web-v1 |
| svc-node-web-v2 | NodePort  | 10.233.26.196 | <none>      | 80:30002/TCP | 7s  | app=node-web-v2 |





# Status

```
remote > kubectl get replicaset -o wide
```

| NAME      | DESIRED | CURRENT | READY | AGE   | CONTAINERS | IMAGES                | SELECTOR        |
|-----------|---------|---------|-------|-------|------------|-----------------------|-----------------|
| rs-web-v1 | 3       | 3       | 3     | 8m55s | node-web   | whatwant/node-web:1.0 | app=node-web-v1 |
| rs-web-v2 | 3       | 3       | 3     | 8m52s | node-web   | whatwant/node-web:2.0 | app=node-web-v2 |

```
remote > kubectl get pods -o wide
```

| NAME            | READY | STATUS  | RESTARTS | AGE   | IP            | NODE    | NOMINATED | NODE | READINESS | GATES |
|-----------------|-------|---------|----------|-------|---------------|---------|-----------|------|-----------|-------|
| rs-web-v1-ghbqs | 1/1   | Running | 0        | 9m24s | 10.233.110.29 | worker1 | <none>    |      | <none>    |       |
| rs-web-v1-m6gtz | 1/1   | Running | 0        | 9m24s | 10.233.103.79 | worker2 | <none>    |      | <none>    |       |
| rs-web-v1-rxpxz | 1/1   | Running | 0        | 9m24s | 10.233.103.80 | worker2 | <none>    |      | <none>    |       |
| rs-web-v2-cmv12 | 1/1   | Running | 0        | 9m21s | 10.233.103.82 | worker2 | <none>    |      | <none>    |       |
| rs-web-v2-k6src | 1/1   | Running | 0        | 9m21s | 10.233.110.30 | worker1 | <none>    |      | <none>    |       |
| rs-web-v2-qsh68 | 1/1   | Running | 0        | 9m21s | 10.233.103.81 | worker2 | <none>    |      | <none>    |       |

```
remote > kubectl get services -o wide
```

| NAME            | TYPE      | CLUSTER-IP    | EXTERNAL-IP | PORT(S)      | AGE | SELECTOR        |
|-----------------|-----------|---------------|-------------|--------------|-----|-----------------|
| kubernetes      | ClusterIP | 10.233.0.1    | <none>      | 443/TCP      | 24d | <none>          |
| svc-node-web-v1 | NodePort  | 10.233.24.241 | <none>      | 80:30001/TCP | 11s | app=node-web-v1 |
| svc-node-web-v2 | NodePort  | 10.233.26.196 | <none>      | 80:30002/TCP | 7s  | app=node-web-v2 |

```
remote > curl -s http://192.168.100.200:30001
```

You've hit rs-web-v1-m6gtz

```
remote > curl -s http://192.168.100.200:30002
```

You've hit rs-web-v2-k6src (Ver2.0)

# Create Ingress (Simple Fanout)

## ingress-node-web.yaml

```
apiVersion: networking.k8s.io/v1
kind: Ingress
```

### metadata:

```
  name: ingress-node-web
```

### annotations:

```
  kubernetes.io/ingress.class: "nginx"
```

```
  nginx.ingress.kubernetes.io/rewrite-target: /
```

### spec:

#### rules:

##### - http:

#### paths:

##### - path: /foo

```
  pathType: Prefix
```

#### backend:

##### service:

```
  name: svc-node-web-v1
```

#### port:

```
  number: 80
```

##### - path: /bar

```
  pathType: Prefix
```

#### backend:

##### service:

```
  name: svc-node-web-v2
```

#### port:

```
  number: 80
```

```
remote > cd kubernetes/05-Ingress-LoadBalancer/hands-on
```

```
remote > kubectl create -f ingress-node-web.yaml
```

```
ingress.networking.k8s.io/ingress-node-web created
```

```
remote > kubectl get ingresses -o wide
```

| NAME             | CLASS  | HOSTS | ADDRESS         | PORTS | AGE |
|------------------|--------|-------|-----------------|-------|-----|
| ingress-node-web | <none> | *     | 192.168.100.201 | 80    | 11s |

```
remote > kubectl get services --namespace ingress-nginx
```

| NAME                               | TYPE      | CLUSTER-IP    | EXTERNAL-IP | PORT(S)                     | AGE |
|------------------------------------|-----------|---------------|-------------|-----------------------------|-----|
| ingress-nginx-controller           | NodePort  | 10.233.63.114 | <none>      | 80:30977/TCP, 443:30168/TCP | 13h |
| ingress-nginx-controller-admission | ClusterIP | 10.233.59.2   | <none>      | 443/TCP                     | 13h |

```
remote > curl -s http://192.168.100.201:30977/foo
```

```
You've hit rs-web-v1-m6gtz
```

```
remote > curl -s http://192.168.100.201:30977/bar
```

```
You've hit rs-web-v2-cmv12 (Ver2.0)
```

# Describe Ingress

```
remote > kubectl describe ingress ingress-node-web
```


```
Name:          ingress-node-web
Labels:        <none>
Namespace:     default
Address:       192.168.100.201
Default backend: default-http-backend:80 (<error: endpoints "default-http-backend" not found>)
```

```
Rules:
  Host      Path  Backends
  ----      -
  *
    /foo    svc-node-web-v1:80 (10.233.103.79:8080,10.233.103.80:8080,10.233.110.29:8080)
    /bar    svc-node-web-v2:80 (10.233.103.81:8080,10.233.103.82:8080,10.233.110.30:8080)
```

```
Annotations:  kubernetes.io/ingress.class: nginx
              nginx.ingress.kubernetes.io/rewrite-target: /
```

```
Events:
  Type    Reason  Age           From          Message
  ----    -
  Normal  Sync    12m (x2 over 13m)  nginx-ingress-controller  Scheduled for sync
```

# Design Thinking





# externalTrafficPolicy

NodePort  
공부하고 이어서 바로 했으면  
더 좋았을 것 같다.

# Ready

- 앞에서 NodePort 실습했던 내용을 활용해서 살펴보도록 하겠다.

```
remote > cd kubernetes/03-RS-DS-JOB-CJ/hands-on
```

```
remote > kubectl create -f rs-web.yaml
```

```
remote > cd kubernetes/04-ClusterIP-NodePort-ExternalName/hands-on
```

```
remote > kubectl create -f svc-web-node.yaml
```

```
remote > kubectl get pods -o wide
```

| NAME         | READY | STATUS  | RESTARTS | AGE   | IP            | NODE    | NOMINATED | NODE | READINESS | GATES |
|--------------|-------|---------|----------|-------|---------------|---------|-----------|------|-----------|-------|
| rs-web-d7sd5 | 1/1   | Running | 0        | 2m33s | 10.233.110.31 | worker1 | <none>    |      | <none>    |       |
| rs-web-h6cpk | 1/1   | Running | 0        | 2m33s | 10.233.103.85 | worker2 | <none>    |      | <none>    |       |
| rs-web-ltscf | 1/1   | Running | 0        | 2m33s | 10.233.103.84 | worker2 | <none>    |      | <none>    |       |

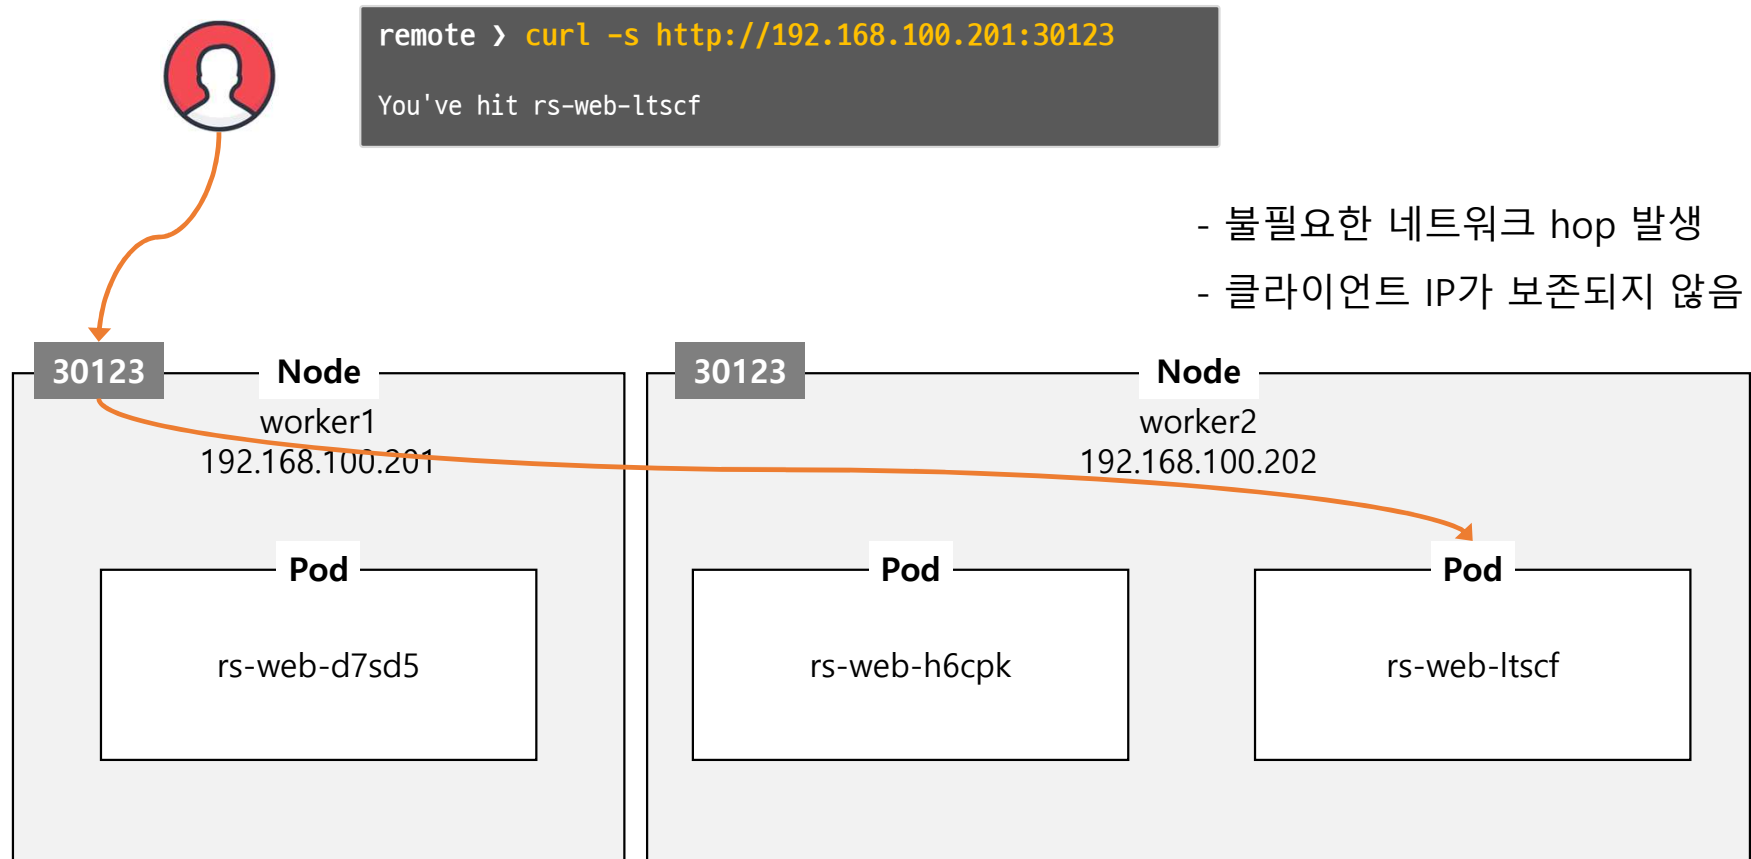
```
remote > kubectl get services -o wide
```

| NAME         | TYPE      | CLUSTER-IP   | EXTERNAL-IP | PORT(S)      | AGE | SELECTOR     |
|--------------|-----------|--------------|-------------|--------------|-----|--------------|
| kubernetes   | ClusterIP | 10.233.0.1   | <none>      | 443/TCP      | 24d | <none>       |
| svc-web-node | NodePort  | 10.233.7.149 | <none>      | 80:30123/TCP | 77s | app=node-web |

```
remote > curl -s http://192.168.100.201:30123
```

```
You've hit rs-web-ltscf
```

# Problem





# externalTrafficPolicy

## svc-web-node-et.yaml

```
apiVersion: v1
kind: Service

metadata:
  name: svc-web-node-et

spec:
  type: NodePort

  ports:
    - port: 80
      targetPort: 8080
      nodePort: 30125

  selector:
    app: node-web

externalTrafficPolicy: Local
```

```
remote > cd kubernetes/05-Ingress-LoadBalancer/hands-on
```

```
remote > kubectl create -f svc-web-node-et.yaml
```

```
service/svc-web-node-et created
```


```
remote > kubectl get services -o wide
```

| NAME            | TYPE      | CLUSTER-IP   | EXTERNAL-IP | PORT(S)      | AGE | SELECTOR     |
|-----------------|-----------|--------------|-------------|--------------|-----|--------------|
| kubernetes      | ClusterIP | 10.233.0.1   | <none>      | 443/TCP      | 24d | <none>       |
| svc-web-node    | NodePort  | 10.233.7.149 | <none>      | 80:30123/TCP | 89m | app=node-web |
| svc-web-node-et | NodePort  | 10.233.6.235 | <none>      | 80:30125/TCP | 27s | app=node-web |

```
remote > curl -s http://192.168.100.201:30125
```

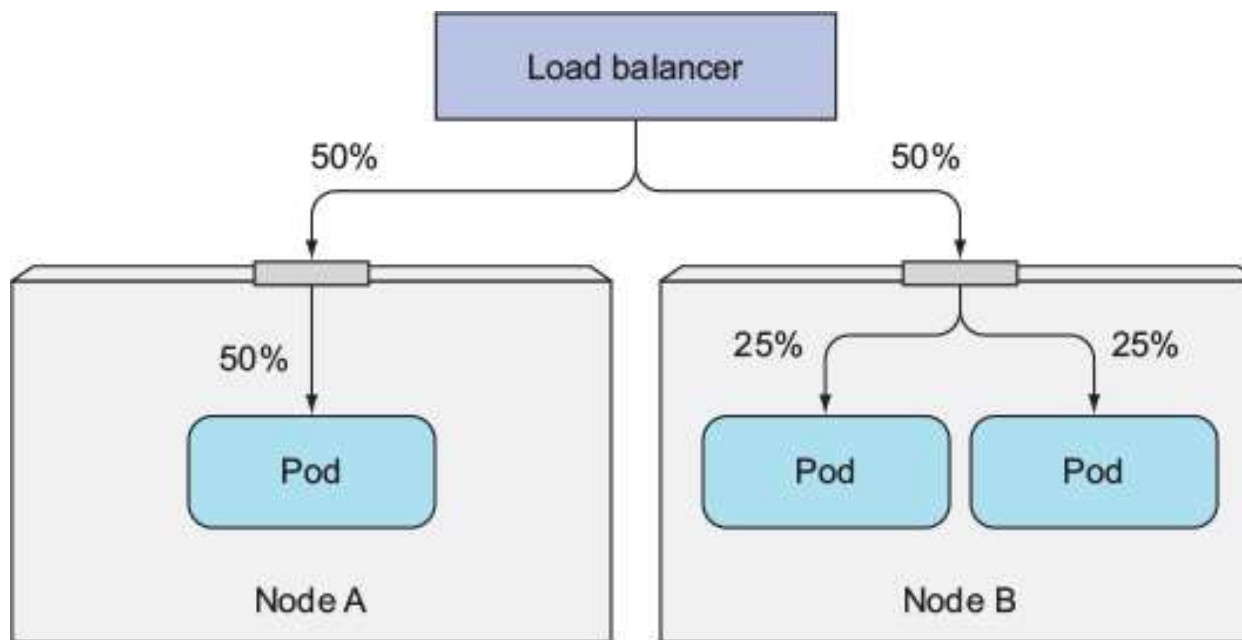
```
You've hit rs-web-d7sd5
```

- externalTrafficPolicy 기본값은 Cluster



# But,

- LoadBalancer를 사용하면서 externalTrafficPolicy를 적용하게 되면, 오히려 균등 배부가 되지 않을 수도 있다.



※ "externalTrafficPolicy: Local" 설정에 따른 문제점 : 균등히 배부되지 않을 수 있음

※ 참고 : <https://livebook.manning.com/book/kubernetes-in-action/chapter-5/230>



**Break**

**돌아오셨으면 채팅창에  
복귀!  
타이핑하기!**


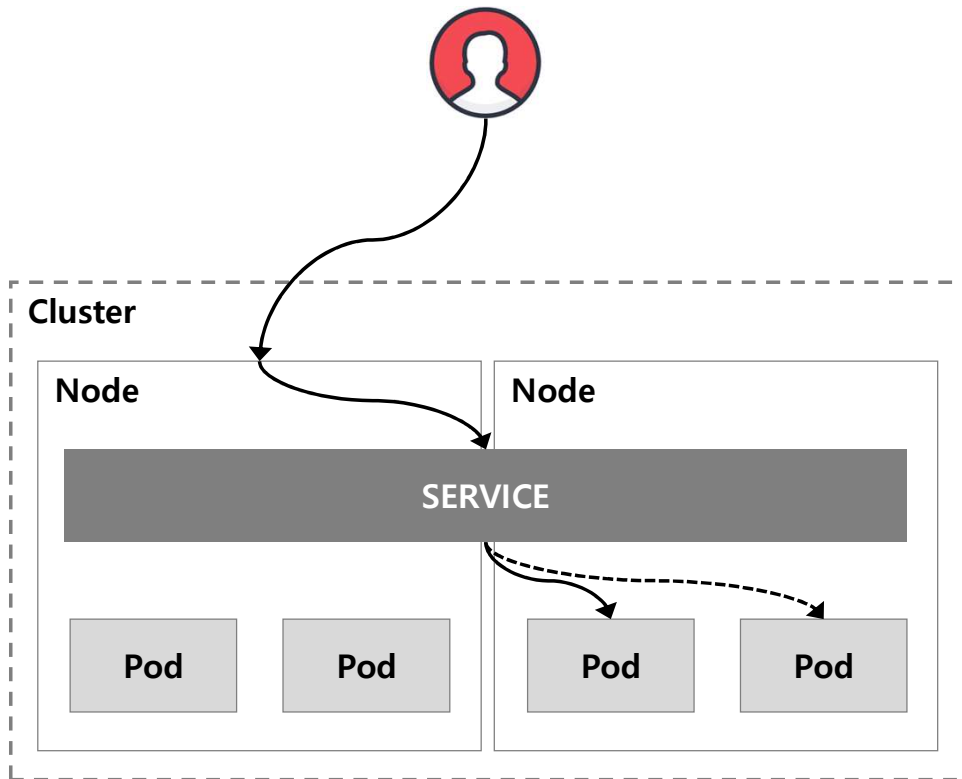


# **Kubernetes**

## **Network - LoadBalancer**

# Why LoadBalancer ?

## NodePort vs LoadBalancer







# **MetalLB**

# **Installation**

# Requirement - 1/3

- MetalLB requires the following to function:

- ① A Kubernetes cluster, running Kubernetes **1.13.0 or later**, that does not already have network load-balancing functionality.
- ② A cluster network configuration that can **coexist with MetalLB**.
- ③ Some **IPv4 addresses** for MetalLB to hand out.
- ④ When using the BGP operating mode, you will need one or more routers capable of speaking BGP.
- ⑤ Traffic on **port 7946** (TCP & UDP) must be allowed between nodes, as required by memberlist.

## ① Kubernetes version

```
remote > kubectl get nodes -o wide
```

| NAME    | STATUS | ROLES                | AGE | VERSION | INTERNAL-IP     | EXTERNAL-IP | OS-IMAGE           | KERNEL-VERSION    | CONTAINER-RUNTIME  |
|---------|--------|----------------------|-----|---------|-----------------|-------------|--------------------|-------------------|--------------------|
| master  | Ready  | control-plane,master | 24d | v1.23.7 | 192.168.100.200 | <none>      | Ubuntu 20.04.4 LTS | 5.4.0-122-generic | containerd://1.6.4 |
| worker1 | Ready  | <none>               | 24d | v1.23.7 | 192.168.100.201 | <none>      | Ubuntu 20.04.4 LTS | 5.4.0-124-generic | containerd://1.6.4 |
| worker2 | Ready  | <none>               | 24d | v1.23.7 | 192.168.100.202 | <none>      | Ubuntu 20.04.4 LTS | 5.4.0-124-generic | containerd://1.6.4 |

Kubernetes 1.13.0 버전 이상이고, 다른 load-balancing 기능을 갖고 있지 않으면 된다.

※ 참고 : <https://metallb.universe.tf/#requirements>

# Requirement - 2/3

## ② Cluster Network

Calico 환경에서는 알려진 문제가 있지만,  
BGP를 사용할 경우에만 해당하는 문제이니 일단 무시하고 진행~

```
remote > kubectl get pods --namespace kube-system
```

| NAME                                     | READY | STATUS  | RESTARTS       | AGE |
|--|-------|---------|----------------|-----|
| calico-kube-controllers-6dd874f784-rjxtz | 1/1   | Running | 14 (4h11m ago) | 24d |
| calico-node-599ck                        | 1/1   | Running | 10 (4h12m ago) | 24d |
| calico-node-qlhvf                        | 1/1   | Running | 10 (4h11m ago) | 24d |
| calico-node-tpwvg                        | 1/1   | Running | 20 (4h11m ago) | 24d |
| coredns-76b4fb4578-kc7vm                 | 1/1   | Running | 10 (4h12m ago) | 24d |
| coredns-76b4fb4578-zbtvb                 | 1/1   | Running | 9 (4h11m ago)  | 24d |
| dns-autoscaler-7979fb6659-p5fpm          | 1/1   | Running | 10 (4h12m ago) | 24d |
| kube-apiserver-master                    | 1/1   | Running | 11 (4h12m ago) | 24d |
| kube-controller-manager-master           | 1/1   | Running | 11 (4h12m ago) | 24d |
| ...                                      |       |         |                |     |

| Network addon | Compatible                          |
|---------------|-------------------------------------|
| Antrea        | Yes (Tested on version 1.4 and 1.5) |
| Calico        | Mostly (see known issues)           |
| Canal         | Yes                                 |
| Cilium        | Yes                                 |
| Flannel       | Yes                                 |
| Kube-ovn      | Yes                                 |
| Kube-router   | Mostly (see known issues)           |
| Weave Net     | Mostly (see known issues)           |

## ③ IPv4 addresses

**DHCP 서버 설정**

공유기의 DHCP 서버 기능을 관리합니다.

DHCP 서버 사용함 : ☒

DHCP IP주소 범위 :  to  (1 ~ 254 사이의 숫자)


MetalLB에 할당할 수 있는 IP들을 확보해야 한다.

※ 참고 : <https://metallb.universe.tf/installation/network-addons/>

| 로컬 클라이언트 상태       |                 |     |
|-------------------|-----------------|-----|
| MAC 주소            | IP 주소           | 장치명 |
| 08:00:27:2D:8C:6E | 192.168.100.102 | vm1 |
| 08:00:27:2D:8C:6F | 192.168.100.107 | vm2 |
| 08:00:27:2D:8C:70 | 192.168.100.101 | vm3 |
| 08:00:27:2D:8C:71 | 192.168.100.200 |     |
| 08:00:27:2D:8C:72 | 192.168.100.201 |     |
| 08:00:27:2D:8C:73 | 192.168.100.202 |     |

## Requirement - 3/3

### ④ When using the BGP(Boarder Gateway Protocol)...



MetalLB는 L2/L3 2가지 방식 중 하나를 선택해서 설치할 수 있다. L3(BGP) 방식을 선택할 경우 추가적으로 고려할 사항들이 있다. 하지만, 여기에서는 일단 L2로 진행할 계획이기에 Skip ...

|           |             |            |             |
|-----------|-------------|------------|-------------|
| <b>L2</b> | Data Link   | Mac 주소 기반  |             |
| <b>L3</b> | Network     | IP 주소 기반   |             |
| <b>L4</b> | Transport   | Port 기반    | TCP, UDP    |
| <b>L7</b> | Application | 요청(URL) 기반 | HTTP, HTTPS |

### ⑤ Traffic on port 7946 (TCP & UDP)

MetalLB가 동작하기 위해서는 Node끼리 port 7946 (TCP/UDP) 통신이 되어야 한다.

보통의 상황에서는 문제가 없겠지만, 방화벽이 있는 환경에서는 해당 port가 열려 있도록 주의해야 한다.

# Preparation

"If you're using kube-proxy in **IPVS** mode, since Kubernetes v1.14.2 you have to enable strict ARP mode"

- 실습환경 구축을 가이드한대로 진행했다고 하면, kube-proxy 기본 설정이 `iptables` 일 것이다.
- 만약, `IPVS` 설정으로 되어 있다면 strictARP 설정을 enable로 설정해줘야 한다.
- kube-proxy 설정 내역을 확인해 보려면 다음과 같이 해보면 된다.

```
remote > kubectl describe configmap kube-proxy --namespace kube-system
```

```
...
iptables:
  masqueradeAll: false
  masqueradeBit: 14
  minSyncPeriod: 0s
  syncPeriod: 30s
ipvs:
  excludeCIDRs: []
  minSyncPeriod: 0s
  scheduler: rr
  strictARP: false
  syncPeriod: 30s
  tcpFinTimeout: 0s
  tcpTimeout: 0s
  udpTimeout: 0s
kind: KubeProxyConfiguration
metricsBindAddress: 127.0.0.1:10249
mode: iptables
...
```

※ 참고 : <https://metallb.universe.tf/installation/>



# Install MetalLB

```
remote > kubectl apply -f https://raw.githubusercontent.com/metallb/metallb/v0.13.4/config/manifests/metallb-native.yaml
```

```
namespace/metallb-system created
customresourcedefinition.apiextensions.k8s.io/addresspools.metallb.io created
customresourcedefinition.apiextensions.k8s.io/bfdprofiles.metallb.io created
customresourcedefinition.apiextensions.k8s.io/bgpadvertisements.metallb.io created
customresourcedefinition.apiextensions.k8s.io/bgppeers.metallb.io created
customresourcedefinition.apiextensions.k8s.io/communities.metallb.io created
customresourcedefinition.apiextensions.k8s.io/ipaddresspools.metallb.io created
customresourcedefinition.apiextensions.k8s.io/l2advertisements.metallb.io created
serviceaccount/controller created
serviceaccount/speaker created
Warning: policy/v1beta1 PodSecurityPolicy is deprecated in v1.21+, unavailable in v1.25+
podsecuritypolicy.policy/controller created
podsecuritypolicy.policy/speaker created
role.rbac.authorization.k8s.io/controller created
role.rbac.authorization.k8s.io/pod-lister created
clusterrole.rbac.authorization.k8s.io/metallb-system:controller created
clusterrole.rbac.authorization.k8s.io/metallb-system:speaker created
rolebinding.rbac.authorization.k8s.io/controller created
rolebinding.rbac.authorization.k8s.io/pod-lister created
clusterrolebinding.rbac.authorization.k8s.io/metallb-system:controller created
clusterrolebinding.rbac.authorization.k8s.io/metallb-system:speaker created
secret/webhook-server-cert created
service/webhook-service created
deployment.apps/controller created
daemonset.apps/speaker created
validatingwebhookconfiguration.admissionregistration.k8s.io/metallb-webhook-configuration created
```

※ 참고 : <https://metallb.universe.tf/installation/#installation-by-manifest>



# Check

remote > **kubectl get namespaces**

| NAME            | STATUS | AGE   |
|-----------------|--------|-------|
| default         | Active | 24d   |
| ingress-nginx   | Active | 16h   |
| kube-node-lease | Active | 24d   |
| kube-public     | Active | 24d   |
| kube-system     | Active | 24d   |
| metallb-system  | Active | 3m12s |

remote > **kubectl get pods -o wide --namespace metallb-system**

| NAME                        | READY | STATUS  | RESTARTS | AGE   | IP              | NODE    | NOMINATED | NODE | READINESS | GATES |
|-----------------------------|-------|---------|----------|-------|-----------------|---------|-----------|------|-----------|-------|
| controller-64cc46b9f9-vxqb2 | 1/1   | Running | 0        | 4m40s | 10.233.103.83   | worker2 | <none>    |      | <none>    |       |
| speaker-dwprs               | 1/1   | Running | 0        | 4m40s | 192.168.100.201 | worker1 | <none>    |      | <none>    |       |
| speaker-n8prv               | 1/1   | Running | 0        | 4m40s | 192.168.100.202 | worker2 | <none>    |      | <none>    |       |
| speaker-prhkx               | 1/1   | Running | 0        | 4m40s | 192.168.100.200 | master  | <none>    |      | <none>    |       |

# Configuration

- L2 / BGP 유형에 따라 설정이 다르지만, 여기에서는 L2 기준으로 진행하겠다.

## metallb-ip.yaml

```
apiVersion: metallb.io/v1beta1
kind: IPAddressPool

metadata:
  name: lb-pool
  namespace: metallb-system

spec:
  addresses:
    - 192.168.100.240-192.168.100.250
```

## metallb-l2.yaml

```
apiVersion: metallb.io/v1beta1
kind: L2Advertisement

metadata:
  name: lb
  namespace: metallb-system

spec:
  ipAddressPools:
    - lb-pool
```

```
remote > cd kubernetes/05-Ingress-LoadBalancer/hands-on
```

```
remote > kubectl create -f metallb-ip.yaml
```

```
ipaddresspool.metallb.io/lb-pool created
```

```
remote > kubectl create -f metallb-l2.yaml
```

```
l2advertisement.metallb.io/lb created
```

※ 참고 : <https://metallb.universe.tf/configuration/#layer-2-configuration>



# Ready

- 앞에서 진행했던 ReplicaSet (Pods) 및 Service를 그대로 사용해서 LoadBalancer를 실습해보도록 하겠다.

```
remote > kubectl get replicaset -o wide
```

| NAME      | DESIRED | CURRENT | READY | AGE   | CONTAINERS | IMAGES                | SELECTOR        |
|-----------|---------|---------|-------|-------|------------|-----------------------|-----------------|
| rs-web-v1 | 3       | 3       | 3     | 8m55s | node-web   | whatwant/node-web:1.0 | app=node-web-v1 |
| rs-web-v2 | 3       | 3       | 3     | 8m52s | node-web   | whatwant/node-web:2.0 | app=node-web-v2 |

```
remote > kubectl get pods -o wide
```

| NAME            | READY | STATUS  | RESTARTS | AGE   | IP            | NODE    | NOMINATED NODE | READINESS GATES |
|-----------------|-------|---------|----------|-------|---------------|---------|----------------|-----------------|
| rs-web-v1-ghbqs | 1/1   | Running | 0        | 9m24s | 10.233.110.29 | worker1 | <none>         | <none>          |
| rs-web-v1-m6gtz | 1/1   | Running | 0        | 9m24s | 10.233.103.79 | worker2 | <none>         | <none>          |
| rs-web-v1-rxpxz | 1/1   | Running | 0        | 9m24s | 10.233.103.80 | worker2 | <none>         | <none>          |
| rs-web-v2-cmvl2 | 1/1   | Running | 0        | 9m21s | 10.233.103.82 | worker2 | <none>         | <none>          |
| rs-web-v2-k6src | 1/1   | Running | 0        | 9m21s | 10.233.110.30 | worker1 | <none>         | <none>          |
| rs-web-v2-qsh68 | 1/1   | Running | 0        | 9m21s | 10.233.103.81 | worker2 | <none>         | <none>          |

```
remote > kubectl get services -o wide
```

| NAME            | TYPE      | CLUSTER-IP    | EXTERNAL-IP | PORT(S)      | AGE | SELECTOR        |
|-----------------|-----------|---------------|-------------|--------------|-----|-----------------|
| kubernetes      | ClusterIP | 10.233.0.1    | <none>      | 443/TCP      | 24d | <none>          |
| svc-node-web-v1 | NodePort  | 10.233.24.241 | <none>      | 80:30001/TCP | 11s | app=node-web-v1 |
| svc-node-web-v2 | NodePort  | 10.233.26.196 | <none>      | 80:30002/TCP | 7s  | app=node-web-v2 |

# Create LoadBalancer

- LoadBalancer 유형의 Service를 생성하는 것은 type 으로 명시만 해주면 된다.

## svc-node-web-lb.yaml

```
apiVersion: v1
kind: Service
metadata:
  name: svc-lb
spec:
  type: LoadBalancer
  ports:
    - name: http
      port: 80
      protocol: TCP
      targetPort: 8080
  selector:
    app: node-web-v1
```

```
remote > cd kubernetes/05-Ingress-LoadBalancer/hands-on
```

```
remote > kubectl create -f svc-node-web-lb.yaml
```

```
service/svc-lb created
```

```
remote > kubectl get services -o wide
```

| NAME            | TYPE         | CLUSTER-IP    | EXTERNAL-IP     | PORT(S)      | AGE   | SELECTOR        |
|-----------------|--------------|---------------|-----------------|--------------|-------|-----------------|
| kubernetes      | ClusterIP    | 10.233.0.1    | <none>          | 443/TCP      | 24d   | <none>          |
| svc-lb          | LoadBalancer | 10.233.46.241 | 192.168.100.240 | 80:31336/TCP | 4s    | app=node-web-v1 |
| svc-node-web-v1 | NodePort     | 10.233.24.241 | <none>          | 80:30001/TCP | 7h54m | app=node-web-v1 |
| svc-node-web-v2 | NodePort     | 10.233.26.196 | <none>          | 80:30002/TCP | 7h54m | app=node-web-v2 |

```
remote > curl -s http://192.168.100.240
```

```
You've hit rs-web-v1-m6gtz
```

# Describe LoadBalancer

```
remote > kubectl describe service svc-lb
```

```
Name: svc-lb
Namespace: default
Labels: <none>
Annotations: <none>
Selector: app=node-web-v1
Type: LoadBalancer
IP Family Policy: SingleStack
IP Families: IPv4
IP: 10.233.46.241
IPs: 10.233.46.241
LoadBalancer Ingress: 192.168.100.240
Port: http 80/TCP
TargetPort: 8080/TCP
NodePort: http 31336/TCP
Endpoints: 10.233.103.79:8080,10.233.103.80:8080,10.233.110.29:8080
Session Affinity: None
External Traffic Policy: Cluster
```

Events:

| Type   | Reason       | Age                   | From              | Message  |
|--------|--------------|-----------------------|-------------------|--|
| ----   | -----        | ----                  | ----              | -----  |
| Normal | IPAllocated  | 3m19s                 | metalb-controller | Assigned IP ["192.168.100.240"]                      |
| Normal | nodeAssigned | 2m2s (x25 over 3m19s) | metalb-speaker    | announcing from node "master" with protocol "layer2" |




**sessionAffinity**  
**(Sticky Session)**



# Problem

Without Session Stickiness



With Session Stickiness



- 앞에서 실습한 LoadBalancer 를 조금 더 살펴보자.

```
remote > curl -s http://192.168.100.240
You've hit rs-web-v1-ghbqs

remote > curl -s http://192.168.100.240
You've hit rs-web-v1-rxpxz

remote > curl -s http://192.168.100.240
You've hit rs-web-v1-m6gtz

remote > curl -s http://192.168.100.240
You've hit rs-web-v1-rxpxz
```

- 접속할 때 마다 매번 바뀌는 Endpoints 를 볼 수 있다 → 만약 게시판과 같은 웹사이트라면 문제가 있을 수 있다!

※ 참고 : <https://www.imperva.com/learn/availability/sticky-session-persistence-and-cookies/>

# sessionAffinity (Sticky Session)

- 세션이 유지되는 동안 접속되는 Pod가 동일하도록 하기 위한 설정을 해보자.

## svc-node-web-lb-sa.yaml

```
apiVersion: v1
kind: Service
metadata:
  name: svc-lb-sa
spec:
  type: LoadBalancer

  ports:
    - name: http
      port: 80
      protocol: TCP
      targetPort: 8080

  selector:
    app: node-web-v1
```

```
sessionAffinity: ClientIP
sessionAffinityConfig:
  clientIP:
    timeoutSeconds: 3600
```

```
remote > cd kubernetes/05-Ingress-LoadBalancer/hands-on
```

```
remote > kubectl create -f svc-node-web-lb-sa.yaml
```

```
service/svc-lb-sa created
```

```
remote > kubectl get services -o wide
```

| NAME       | TYPE         | CLUSTER-IP    | EXTERNAL-IP     | PORT(S)      | AGE | SELECTOR        |
|------------|--------------|---------------|-----------------|--------------|-----|-----------------|
| kubernetes | ClusterIP    | 10.233.0.1    | <none>          | 443/TCP      | 24d | <none>          |
| svc-lb     | LoadBalancer | 10.233.46.241 | 192.168.100.240 | 80:31336/TCP | 49m | app=node-web-v1 |
| svc-lb-sa  | LoadBalancer | 10.233.20.177 | 192.168.100.241 | 80:30566/TCP | 25s | app=node-web-v1 |

```
remote > curl -s http://192.168.100.241
```

```
You've hit rs-web-v1-ghbqs
```

```
remote > curl -s http://192.168.100.241
```

```
You've hit rs-web-v1-ghbqs
```

```
remote > curl -s http://192.168.100.241
```

```
You've hit rs-web-v1-ghbqs
```

- `sessionAffinityConfig` 기본값은 10,800sec 이다. 즉, 지정하지 않아도 되는 설정이지만 공부를 위해서 추가해보았다.



**<https://kahoot.it/>**



# 자습 (복습)

▷ Ingress

▷ LoadBalancer